1 Table of Contents

1 Table of Contents 1

2 Purpose of this RFP 5

3 Background and Business Objectives 5

4 Expected Benefits 6

5 Technology and Service Provision 6

5.1 Vendor Submission Guidelines for LOTs 7

5.2 Clarification and Q&A Process 7

5.2.1 Submission Requirements 8

5.2.2 Submission and Response Timeline 8

5.2.3 Rules of Engagement 8

6 LOT A – Enterprise AI/RAG Control Plane Components 9

6.1 LOT A Governance Authority 11

6.2 Mandatory platform components 11

6.3 Optional platform component for Agentic 13

6.4 Other Optional platform components 14

6.5 LOT A / LOT B Responsibility Allocation 15

6.6 Group AI Infrastructure Alignment (Strategic Context) 16

7 LOT Bx – Use-case work packages 16

7.1 Unified Vendor Obligations for LOT Bx – Use-Case Delivery 17

7.1.1 Mandatory Consumption of the Enterprise AI/RAG Control Plane 17

7.1.2 Functional & Technical Design Obligations 18

7.1.3 Governance, Traceability & Explainability 18

7.1.4 Security, GDPR, EU AI Act & DORA Compliance 18

7.1.5 Deployment, CI/CD & Environment Obligations 19

7.1.6 Testing, Performance & Drift Management 19

7.1.7 Documentation, Training & Operational Readiness 19

7.1.8 Prohibition of Ungoverned Vendor-Side Infrastructure 19

7.2 LOT B1 – P&C Corporate Underwriting – Use Case Requirements 20

7.3 LOT B2 – CX Contact Centre – Use Case Requirements 20

7.4 LOT B3 – Group Life & Health – Use Case Requirements 20
7.5 LOT B4 – Life & Health Claims – Use Case Requirements 20
7.6 LOT B5 – Life & Health Underwriting – Use Case Requirements 21
7.7 LOT B6 – P&C Motor Claims – Use Case Requirements 21
7.8 LOT B7 – Technology – Use Case Requirements 21
7.9 Certified External Platform (CEP) Pathway 22
7.9.1 Eligibility Criteria for CEP Designation 22
7.9.2 Mandatory CEP Integration Obligations 22
7.9.3 LOT A Vendor Responsibilities for CEP Governance 23
7.9.4 CEP Evaluation and Scoring 23
8 Information Security & DORA Requirements 24
9 Data Privacy & GDPR Requirements 25
10 EU AI Act Compliance Requirements 27
11 Technical Evaluation Criteria 28
11.1 Evaluation Criteria for LOT A (Control Plane Vendor) 29
11.2 Evaluation Criteria for LOT Bx (Use-Case Vendors) 29
12 Legal & Contractual Requirements 31
12.1 Contracting Structure 31
12.2 Data Protection & Regulatory Compliance 31
12.3 Intellectual Property Rights 32
12.4 Confidentiality 32
12.5 Sub-contracting & Sub-processors 32
12.6 Security & Audit Rights- Regulatory Cooperation 33
12.7 Service Levels & Service Credits 33
12.8 Change Control 33
12.9 Exit, Transition & Portability 33
12.10 Indemnities & Liability 34
12.11 Insurance Requirements 34
12.12 Additional Requirements 34
13 Annexes 36
13.1 Annex A – Control Plane Requirements (LOT A only) 36
13.1.1 Purpose of the Control Plane 36

13.1.2 Retrieval & Vector Search Layer 36
13.1.3 Ingestion, Pre-Processing & Embedding Services 37
13.1.4 LLM Gateway & Model Routing 37
13.1.5 Prompt Management & Governance Layer 37
13.1.6 Governance, Auditability & Explainability 38
13.1.7 Monitoring, Observability & AI-Ops 38
13.1.8 Environment Management, CI/CD & Operational Readiness 38
13.1.9 Vendor Onboarding & Multi-Vendor Integration Requirements 38
13.1.10 Operational & Service Management Requirements 39
13.1.11 Data & Content Lifecycle Governance 39
13.1.12 Mandatory Deliverables for LOT A Vendor 39
13.1.13 Optional Value-Add Platform Capabilities 40
13.2 Annex B – API Specifications 41
13.2.1 Purpose 41
13.2.2 Retrieval API (LOT A Must Deliver) Capabilities & LOT B Obligations 41
13.2.3 Ingestion & Embedding API Capabilities & Behaviors & LOT B Obligations 41
13.2.4 LLM Gateway API Capabilities & Outputs & LOT B Obligations 42
13.2.5 Prompt Management API Capabilities & LOT B Obligations 42
13.2.6 Audit & Explainability API Capabilities & LOT B Obligations 42
13.2.7 Observability & Monitoring API Capabilities & LOT A/B Obligations 42
13.2.8 API Security Requirements 43
13.2.9 Performance & Reliability 43
13.2.10 API Versioning & Lifecycle 43
13.2.11 Sandbox & Testing Toolkit (LOT A Must Deliver) 43
13.2.12 Schema Finalization 43
13.2.13 Optional Agent APIs (Required if Component H is Proposed) 43
13.3 Annex C – Use-Case Definitions 45
13.3.1 Annex C1 – Executive list of use-cases 45
13.3.2 Annex C2 – Detailed list of use-cases 46
13.4 Annex E – Glossary 49
13.5 Annex F – Compliance Framework (GDPR, EU AI Act, DORA) 51
13.5.1 Annex F1 – GDPR Compliance Pack 51

13.5.2 Annex F2 – EU AI Act Compliance Pack 51
13.5.3 Annex F3 – DORA ICT Risk & Operational Resilience Pack 52
13.6 Annex H – Architecture & Integration Requirements 53
13.6.1 High-Level Enterprise AI/RAG Control Plane Architecture 53
13.6.2 Integration Principles for LOT Bx Vendors 53
13.6.3 Environment Model & Operational Architecture 54
13.6.4 Security Architecture Requirements (Extended) 54
13.6.5 Data Flow, Lifecycle & Governance Framework 54
13.6.6 Vendor Onboarding & Compliance Sequence 55
13.6.7 Performance, Scalability & Capacity Requirements 55
13.6.8 Logging, Traceability & Explainability Framework 55
13.6.9 Prohibited Architectural Patterns 55
13.7 Annex I – Vendor Compliance Checklists (LOT A & LOT B) 56
13.7.1 Annex I.1 - General instructions applicable to all LOTs 56
13.7.2 Annex I.2 - Instructions for LOT A Bidders (Control Plane Vendor) 57
13.7.3 Annex I.3 - Instructions for LOT B Bidders (Use-Case Vendors) 58
13.7.4 Annex I.4 - LOT A Vendor Compliance Checklist 59
13.7.5 Annex I.5 - LOT B Vendor Compliance Checklist 60
13.8 Annex J: Ethniki Insurance Technology Landscape 62
13.9 Annex K – Use-Case Delivery Playbook (Tentative Framework for Bidder Consideration) 63
13.9.1 Delivery Phases 63
13.10 Annex L – LOT A Control Plane Delivery Playbook (Mandatory) 64
13.10.1 Delivery Phases 64
13.11 Annex M – Group AI Infrastructure: Strategic Context and Evaluative Preference 65
13.11.1 Background 65
13.11.2 The Piraeus Group AI Foundry Architecture 65
13.11.3 Why Group Alignment Is Strategically Attractive to Ethniki 66
13.11.4 Status of Due Diligence and Implications for This Procurement 67
13.11.5 Guidance for LOT A Vendors 67
13.12 Annex N – Q&A Submission Template 67

2 Purpose of this RFP

The purpose of this Request for Proposal (RFP) is to invite qualified vendors to submit proposals to participate in the delivery of Ethniki Insurance's enterprise-wide Artificial Intelligence (AI) and Intelligent Document Processing (IDP) transformation program.

This RFP establishes a two-layer sourcing framework:

- LOT A – Provision, implementation and operation of the Enterprise AI/Retrieval-Augmented Generation (RAG) Control Plane.
- LOTs B1–B7 – Delivery of AI and IDP use-case work packages by Line of Business (LoB), fully leveraging the Control Plane.

Ethniki Insurance seeks a secure, governed, scalable and multi-vendor delivery ecosystem aligned with GDPR, EU AI Act, and DORA obligations.

3 Background and Business Objectives

Aligned with our strategic goal of fostering innovation and pursuing ongoing enhancement in customer service and operational effectiveness, Ethniki Insurance is committed to substantially elevating customer experience through all interaction channels. The integration of Artificial Intelligence technologies will expand the scope and sophistication of our business operations, which we anticipate will drive improvements in customer experience through reduced turnaround times for both primary and secondary operational processes. This progress will influence performance across customer service (calls and tickets), underwriting, and claims management in all lines of business.

Ethniki Insurance is currently positioned at a pivotal moment. Within the Operations division, two AI and automation-driven solutions have been successfully deployed to production, resulting in marked improvements in turnaround times for these areas. Collectively, these initiatives have produced notable gains in efficiency, cost savings, and enhancements in the quality of front-line services. These achievements underscore the value proposition of AI and automation within our organization: when leveraged for repeatable, data-intensive processes, such technologies consistently generate sustainable operational advantages. Ethniki intends to significantly broaden the use of similar solutions across various scenarios to drive further advancement in our value models. This RFP covers the analysis, design, implementation, deployment, and governance of these initiatives.

Our organization is committed to implementing AI-driven automation and decision-support solutions across all principal business functions, including Underwriting, Claims, Customer Experience, Group Life & Health, and Motor.

Program objectives:

- Improve decision quality and processing speed.
- Reduce manual workload and operational inefficiencies.
- Provide real-time support for customer-facing operations.
- Enforce enterprise-grade governance, risk and compliance controls.
- Build reusable enterprise AI assets and accelerate delivery cycles.

4 Expected Benefits

Our goal with this overall initiative is to optimize our front-line performance and experience and increase our internal operating efficiency, thus positioning Ethniki Insurance as a leader in digital innovation within the local insurance market through increased customer satisfaction, engagement and retention. Specifically:

- Increase our operating efficiency of underwriting and claims processes by decreasing drastically the working time on several departmental streams and procedures and simultaneously ensure the correctness and accuracy of the tasks that will be replaced by AI.
- Optimize our service level by automating simple and repetitive processes, retrieve and populate (on CRM) policy information, assist call center operators and ticket handlers etc.

This way, we intend to strengthen our value proposition and become better differentiated and better preferred in the market through better and more comprehensive customer service, improved customer experience and satisfaction, increase company's attractiveness among independent distribution channels and increased customer loyalty.

Expected benefits include:

- significant% reduction in manual processing workload.
- Enhanced underwriting quality through structured extraction and automated pricing support.
- Faster, more consistent claims processing.
- Improved customer experience via real-time AI assistance.
- Lower TCO via shared platform services.
- Improved compliance posture through embedded safeguards and governance.

5 Technology and Service Provision

The purpose of the procurement that forms the objective of this RFP is project governance services, analysis, solutions/models design, testing, training and implementation services to roll-out AI/IDP and Automation solutions across several functions of operation that meet or exceed the requirements of this RFP (LOTs B1-B7) while maintaining the necessary degree of AI/IDP and Automations Governance through and enterprise wide control plane (LOT A).

Candidate vendors need to ensure their proposal(s) follows the turn-key model, and to ensure that it includes all required aspects of the technology, business capabilities and services required to build and later operate the proposed solution(s), extend and anticipate synergies for cross-departmental established solutions, including software licenses.

Individual use cases are divided into LOTs with the intention being - among other things - to accelerate overall roll-out through parallelism of the initiatives involved. Prospective vendors will be required to submit separate technical and financial proposals for each LOT they wish to be considered for and need to ensure that their offers are not contingent on the result of any other LOT. Should a single Vendor be awarded more than one engagement, the expectation will then be that the different LOTs will be executed concurrently.

5.1 Vendor Submission Guidelines for LOTs

Vendors may submit proposals for any LOT or combination of LOTs (LOT A, LOTs B1 through B7). **However, each LOT will be evaluated independently, therefore, pricing or solution assumptions for one LOT should not be contingent on the outcome of any other LOT.**

The following examples constitute valid submissions:

- Vendor X submits a bid for LOT A only.
- Vendor X submits a bid for LOT B2 only.
- Vendor X submits separate bids for LOT A and LOT B2.
- Vendor X submits a bid for LOT B4, addressing only use cases LHC-3 and LHC-1 (sub-LOT proposal).

Where a Vendor's capability is concentrated in a subset of use cases within a LOT, Ethniki may, at its discretion, consider proposals addressing individual use cases within a LOT rather than the full LOT scope. Vendors wishing to propose on a sub-LOT basis must clearly identify the specific use cases they are addressing, confirm that their proposal satisfies all obligations in Section 7.1 for those use cases, and price each use case independently in accordance with Annex D. Sub-LOT proposals will be evaluated on the same criteria as full LOT proposals for the use cases covered.

Each Vendor Technical offer submission is required to follow a completed compliance checklist – see Sections under 13.8. This applies to bids for each of any of LOT A and LOT Bx.

5.2 Clarification and Q&A Process

Ethniki Asfalistiki operates a structured written Q&A process for this RFP. All clarification requests must be submitted in writing using the Q&A Submission Template provided in Annex N. Verbal, informal, or unstructured questions will not receive official responses and will not be binding on Ethniki. All questions and official responses will be shared simultaneously with all registered bidders to preserve procurement impartiality. Ethniki reserves the right to amend the RFP in

response to clarification questions; any such amendments will be issued as a numbered addendum and communicated to all registered bidders.

5.2.1 Submission Requirements

Each question must be submitted as a separate row in the Q&A template (Annex N). Bundled or multi-part questions within a single row will be returned for resubmission. The following fields are mandatory for every question:

- **LOT Reference.** The LOT to which the question relates (LOT A; LOT B1 through LOT B7; or General / All LOTs).
- **Section Reference.** The RFP section number to which the question relates (e.g. Section 7.9.2, Annex H.9). Questions that cannot be anchored to a specific section will be returned for clarification.
- **Use Case Reference (if applicable).** The use case code(s) to which the question relates (e.g. PCM-1, LHC-3). Leave blank if the question applies to all use cases within the LOT or is platform-level.
- **Question Type.** Classify the question as one of: Technical / Architectural; Commercial / Pricing; Legal / Contractual; Compliance / Regulatory; or Process / Procedural.
- **Requirement Text Referenced.** A verbatim extract of the specific requirement sentence or bullet point to which the question relates. This field is mandatory and allows Ethniki to locate the precise obligation being questioned without ambiguity.
- **Question.** A concise, self-contained question in English. The question must be specific and unambiguous. Questions seeking general information about Ethniki's preferences outside the RFP scope will not be answered.
- **Proposed Clarification or Interpretation (optional).** Where the Vendor has a proposed reading of the requirement, this should be stated. Ethniki will confirm, correct, or qualify the proposed interpretation in its response. This field accelerates response turnaround and is strongly encouraged.

5.2.2 Submission and Response Timeline

The Q&A timeline will be communicated by Ethniki’s Procurement function at the time of RFP issuance. All submissions must be sent to the designated procurement contact address provided in the covering communication. Questions received after the Q&A deadline may not receive an official response. Responses will be distributed simultaneously to all registered bidders in a consolidated Q&A log, with individual vendor identities redacted.

5.2.3 Rules of Engagement

- Only written responses issued by Ethniki’s Procurement function constitute official clarifications. No other communication — verbal, informal, or otherwise — is binding on Ethniki or modifies the RFP.

- Ethniki reserves the right to decline to answer questions that are commercially sensitive, outside the scope of the RFP, or that would confer an unfair advantage on any individual bidder.
- Where a clarification response materially changes a requirement, Ethniki will issue a formal RFP addendum and, if appropriate, extend the bid submission deadline.
- Vendors must not contact Ethniki business or technical staff directly with questions relating to this RFP. All contact must be channelled through the designated Procurement contact. Breach of this rule may result in disqualification.

6 LOT A – Enterprise AI/RAG Control Plane Components

The LOT A vendor shall design, implement and operate a central AI Control Plane providing all core enterprise AI technical capabilities. The platform will enable Ethniki to securely operate AI across critical business functions. As we expand the use of AI to support underwriting, claims, operations, customer service, and internal knowledge processes, the selected solution must provide a controlled, explainable, and auditable AI foundation that meets our regulatory, security, and performance requirements.

Segregated, traditional LLM deployments are insufficient for enterprise use because they do not inherently access internal systems, cannot verify the accuracy of responses, and pose unacceptable risks in regulated environments. A RAG-based approach resolves these challenges by grounding AI outputs directly on authoritative enterprise data and content. This ensures that answers are accurate, traceable, and aligned with internal policies and regulatory expectations. The control plane approach will need to provide:

- Robust ingestion and vectorization capabilities for all enterprise knowledge sources (documents, databases, procedures, policy wordings, product rules).
- A scalable vector store and retrieval orchestration layer that ensures high relevance, low latency, and full lineage from source to generated answer.
- Secure integration with leading Large Language Models (for Vendors of LOT Bx), with enforced grounding, hallucination controls, and built-in output safety.
- Enterprise-grade governance, security, and compliance, including GDPR, DORA, EU AI Act alignment, auditability, and human-in-the-loop mechanisms.
- Open, API-first integration with existing systems, enabling domain-specific assistants that support our knowledge workers and operational teams.
- Full observability and monitoring across retrieval quality, model usage, accuracy metrics, and operational performance.

Based on the above, the Control Plane will be the **mandatory enterprise AI backbone** that *all* use-case Vendors will need to onboard to. The platform will standardize AI across all use-case Vendors for:

- Vector repository, retrieval, orchestration

- Model gateway / routing
- Data ingestion
- Prompt management
- Cost monitoring & Benefit monitoring
- Observability
- API gateway for interactions

Without the control plane, Ethniki Insurance's AI implementation would become fragmented, ungoverned and unscalable. It is important to understand that Vendors are expected to demonstrate their ability to deliver a complete, production-ready platform, not just individual components. The successful solution must support our strategic goals of improving operational efficiency, accelerating decision-making, ensuring consistency in knowledge application, and enabling responsible use of AI at scale across the organization. In addition, the proposed solution (incl. services) is required to efficiently and effectively support the Vendors of LOT Bx, to onboard their individual solutions to this environment, through sufficient support / documentation / governance mechanisms. It follows that Vendors will need to pair their technology proposal with adequate professional services (incl. QA) to support the transformation program.

The LOT A Vendor shall provide comprehensive enablement, onboarding, governance, testing support, and on-demand technical assistance to all LOT B Vendors.

The evaluation of LOT A will focus on architectural robustness, security, retrieval accuracy, governance maturity, scalability, and total cost of ownership clarity and size. Vendors must articulate how their solution delivers measurable business value while maintaining enterprise safety and compliance.

The awarded partner will play a central role in our AI transformation roadmap and must therefore demonstrate both technical excellence and a strong commitment to long-term partnership.

The following sections, 6.2 and 6.4, describe the functional, technical, operational, and compliance requirements for LOT A (and the corresponding LOT B integration obligations). Mandatory components form the evaluation baseline. Optional components may differentiate vendors but are not required for compliance. Therefore, these sections are <u>solution-defining</u>. They:

- Outline what the solution (professional services and technological elements) needs to accomplish.
- Define technical obligations, interfaces, architecture, constraints, and quality levels.
- Form the basis for technical proposal evaluation.
- Allow candidate Vendors to design, price, and architect their solution.

Section 13.1 (Annex A) lists the documents, artefacts, configurations, datasets, reports, and evidence that the LOT A vendor must produce and deliver throughout the project. Therefore, Annex A is deliverables-defining. It:

- Stipulates for the candidate Vendors what they must provide as outputs.
- Defines contractual deliverables tied to acceptance, onboarding, audit, compliance, and ongoing operations.
- Serves as the basis for contract enforcement and acceptance criteria.
- Allows project management, governance, and steering committees to track progress.

LOT Bx Vendors are advised to refer to these sections to understand what the LOT A vendor must deliver to them as onboarding support.

6.1 LOT A Governance Authority

The LOT A Vendor, acting under Ethniki's governance framework, is responsible for validating and certifying that each LOT B deliverable complies with the architectural, security, retrieval, grounding, observability, logging and compliance requirements of the Enterprise AI/RAG Control Plane.

LOT A is authorized to:

- request remediation from LOT B Vendors,
- block deployment of non-compliant components,
- escalate violations to Ethniki's program governance structures, and
- certify use cases only once all mandatory criteria are satisfied.

No LOT B deliverable may progress to PRE-PROD or PROD without LOT A's technical certification.

LOT B vendors cannot deliver anything unless LOT A onboards, supports & guides them; LOT A must support all LOT B use cases and vendor teams. In addition, any Control Plane that may proposed will not be plug-and-play and will require significant professional services to set-up and operate during on-boarding (including QA, Vendor support etc.). Based on this, LOT A bidders should ensure that the enablement, support, and governance effort required from the LOT A Vendor in relation to all LOT B Vendors is a core part of their solution offering for LOT A. Bidders must allocate sufficient professional services resources to ensure timely onboarding, continuous support, technical certification, and compliance alignment for all LOT B work packages.

6.2 Mandatory platform components

These capabilities **are essential** to deliver a safe, governed, scalable and production-ready enterprise AI platform.

A. Retrieval & Vector Search Layer: Provides governed, enterprise-grade semantic search so all AI outputs can be reliably grounded on approved content.

- Scalable vector DB (dense, sparse, hybrid retrieval).
- Metadata filtering, role-based access propagation.
- Index management and automated retention control.

B. Ingestion & Embedding Services: Transforms raw documents into structured, machine-understandable content, ensuring consistency and accuracy in retrieval.

- Multi-format ingestion (PDF, email, images, XLSX and so on).
- OCR/ICR, form parsing and semantic chunking.
- Automated embedding generation and re-embedding workflows.

C. LLM Gateway & Routing: centralizes access to all LLMs, enforcing safety, compliance, routing, and logging for controlled enterprise AI use.

- Central gateway for multiple LLM families.
- Routing based on cost, latency, sensitivity or risk.
- Enforced grounding and safe-completion.

D. Prompt Management Layer: Ensures prompts are versioned, governed and secure, preventing inconsistent behaviours across agents and workflows.

- Version-controlled prompt library.
- Guardrails and prohibited-content filters.
- Parameterized prompts for retrieval augmentation.

E. Governance, Auditability & Explainability: Captures full traceability for every AI interaction, enabling regulatory compliance and trustworthy operations.

- Logging of prompts, retrieved context, and generated outputs.
- Mandatory source citation for generated answers.
- Explainability tooling for high-risk workflows.

F. Monitoring, Observability & AI-Ops: Provides continuous insight into quality, drift and performance, enabling proactive issue detection and platform stability.

- Dashboards for retrieval quality and hallucination metrics.
- Drift monitoring for models and embeddings.
- SIEM integration for anomaly and security alerts.
- Observability on Ethniki's Dynatrace stack.

G. Platform Operations: Delivers resilient, SLA-backed environments ensuring continuous availability and disaster readiness for all AI workloads.

- 24/7 operational support with defined SLAs.
- RTO ≤ 4 hours, RPO ≤ 0 minutes.
- CI/CD pipelines available to all LOT B vendors.

6.3 Optional platform component for Agentic

H. Agent Orchestration & Execution Layer: A governed runtime enabling multi-step, tool-calling AI agents to operate within the Control Plane boundary. **This capability is designated optional** for the initial procurement but is expected to become mandatory in a subsequent RFP version as Ethniki's agentic AI programme matures. LOT A Vendors are encouraged to propose Component H and will be evaluated on the maturity, governance depth, and roadmap credibility of their agentic offering. Where Component H is proposed, it must satisfy all six requirements below:

H.1 – Managed Agent Runtime. A production-grade runtime for executing multi-step AI agents, supporting sequential, parallel, and conditional execution patterns; stateful multi-turn agent execution with configurable step limits and timeout controls; agent-to-agent delegation and result aggregation; graceful error handling, retry logic, and deterministic fallback paths; and CI/CD pipeline compatibility for agent deployment across DEV/UAT/PROD environments. The vendor may propose any leading orchestration framework as the underlying engine, provided the LOT A governance wrapper satisfies requirements H.2 through H.6.

H.2 – Governed Tool Registry and Execution Sandbox. A centralised registry of all tools and external APIs available for agent use, with mandatory metadata (tool name, owning system, data classification, permitted callers, approval status); a scoped permission model blocking unapproved tool calls; sandboxed execution preventing direct filesystem or network access beyond registered integrations; full execution logging of every tool invocation; tool-level rate limiting and circuit breakers; and routing of all agent LLM calls through the existing LOT A LLM Gateway.

H.3 – Agent Session and State Management. Short-term within-session memory maintaining context and intermediate results throughout an agent's execution lifecycle; long-term cross-session episodic memory for agent-generated artefacts with configurable retention and deletion policies aligned to GDPR data minimisation requirements; state isolation between concurrent agent sessions; and secure encryption of agent state at rest and in transit.

H.4 – Agent Identity and Scoped RBAC. A distinct registered identity (service principal or equivalent) per deployed agent with a defined permission scope; least-privilege

enforcement at registration; clear distinction in audit logs between actions taken by human users and by AI agents; credential rotation policy consistent with enterprise security standards; and alignment with the enterprise identity model (e.g. Active Directory / Entra ID or equivalent) for consistent RBAC enforcement.

H.5 – Human-in-the-Loop (HITL) as Orchestration Primitive. Agents shall support pausing execution at defined HITL checkpoints, routing a structured decision request to a designated human reviewer, and resuming with the reviewer's input incorporated. HITL requests shall include the agent's current state summary, the specific decision required, available options, and consequences. Time-bounded escalation paths shall be configurable where a reviewer does not respond within a defined SLA window. All HITL events shall be fully logged, and a HITL dashboard shall enable supervisors to view pending agent interrupts and review the history of HITL interactions.

H.6 – Agent Observability and Audit Trail. Complete agent execution traces capturing every step, tool call, LLM call, state transition, HITL event, and final output with full lineage; agent-specific metrics (steps per execution, tool call counts, LLM token consumption, HITL frequency, success/failure rates, execution latency); integration with Dynatrace as specified in Section 6.2(F); an immutable audit log for all agent actions involving policyholder data, financial transactions, or regulatory decisions; and EU AI Act transparency requirements for high-risk use cases, with agent decision logs including the basis for each significant decision in human-readable form suitable for regulatory inspection.

Where Component H is proposed, the three corresponding Agent APIs defined in the optional section of Annex B (Agent Orchestration API, Tool Registry API, Agent Audit API) must also be delivered. LOT B Vendors may consume Component H capabilities where offered by the LOT A Vendor, for use cases requiring agentic workflows. LOT B Vendors proposing solutions for use cases identified as having agentic characteristics (e.g. PCM-1, CX-4, LHC-2, CPCU-4, TECH-1, TECH-2) are encouraged to describe their intended approach to agent orchestration in their technical proposal, whether or not the LOT A Vendor proposes Component H.

6.4 Other Optional platform components

These capabilities are **not required** for compliance but may provide functional or performance advantages. They will be evaluated as value-add items only.

- **Advanced Vector Indexing** – Optional accelerators improving retrieval speed and performance at large scale.
- **Multilingual & Translation Layer** – Enables seamless EN/GR content interoperability and multilingual semantic search.
- **Document Similarity & Analytics** – Provides clustering, topic detection and similarity scoring for analytical enrichment.

- **Knowledge Quality Scoring** – Detects content gaps, readability issues and structural inconsistencies in knowledge sources.
- **GenAI Workbench** – Safe experimentation sandbox for developing new prompts, workflows and prototypes without impacting production.

6.5 LOT A / LOT B Responsibility Allocation

To support clarity during bidding and evaluation, Ethniki provides the following indicative RACI matrix, illustrating the expected allocation of responsibilities between the LOT A Control Plane Vendor, the LOT B Use-Case Vendors, and Ethniki Insurance.

This matrix is not definitive and will be validated, refined, and contractually confirmed during the project initiation phase with the selected LOT A Vendor. LOT B Vendors should nonetheless assume that the responsibilities below reflect Ethniki's **intended operating model** and must structure their proposals accordingly. Note, solution certification blocking rights, remediation, and escalation authority remains with LOT A Vendor as per Section 6.1.

Responsibility Area	LOT A Vendor	LOT B Vendor	Ethniki
Platform Architecture & Design	A/R	C	C
Control Plane Build & Operation	A/R	I	C
Environment Provisioning & Access Control	A/R	I	C
Ingestion, Embeddings & Retrieval Infrastructure	A/R	C	C
Prompt Governance & Model Gateway Policies	A	R (usage)	C
Reference Patterns, SDKs & Tooling for LOT B	A/R	C	C
Use-Case Functional Design	C	A/R	R
Use-Case Technical Design	R (alignment)	A/R	C
Security, GDPR, EU AI Act & DORA Baselines	A	A/R (implementation)	R
Test Harness, Certification Criteria & Performance Baselines	A/R	C	C
Use-Case Testing (Functional/Performance/Security)	C	A/R	R
Technical Certification for Deployment	A	R (evidence)	C
CI/CD Deployment & Environment Promotion	A/R	R (artefacts)	C
Operational Monitoring, Logging & Drift Detection	A	R (use-case)	C
Incident Response & Troubleshooting	R (platform)	R (use-case)	A
Change Management & Release Coordination	A	R	R
Documentation & Training	R (platform)	R (use-case)	C

R = Responsible | A = Accountable | C = Consulted | I = Informed

6.6 Group AI Infrastructure Alignment (Strategic Context)

Ethniki Asfalistiki is a member of the Piraeus Bank Group. Piraeus Bank has established and operationalised an enterprise AI Control Plane on Microsoft Azure AI Foundry, covering the equivalent platform components to those specified in Sections 6.2 and 6.3 of this RFP, and currently supporting over 45 AI use cases in production or active development. Ethniki is actively exploring the strategic and commercial case for aligning its LOT A Control Plane implementation with the Piraeus Group AI infrastructure standard, with a view to establishing Group-wide economies of scale in platform governance, regulatory compliance, security operations, and vendor management.

This exploration is ongoing and does not constitute a mandatory architecture requirement for this procurement. The RFP remains open to any qualified LOT A vendor capable of delivering the platform capabilities specified in Sections 6.2 and 6.3, regardless of technology stack. However, Ethniki considers alignment with, or demonstrated compatibility with, a cloud-native Azure AI Foundry-based Group infrastructure blueprint to be a strategically advantageous characteristic in a LOT A proposal. Vendors with relevant Azure AI Foundry implementation experience and a credible plan for Group infrastructure interoperability are encouraged to articulate this capability in their proposals; it will be taken into account within the Vendor Innovation & Delivery Model evaluation criterion (Section 11.1). The strategic context and rationale for this preference are set out in full in Annex M.

7 LOT Bx – Use-case work packages

Since Ethniki plans to procure multiple AI/IDP or Automation use cases across different business domains, the required architectural and procurement model is:

- **LOT A:** The Enterprise AI/RAG Control Plane (Mandatory Foundation)
- **LOTs B1-B7**: The Use-Case Implementation Work Packages; 20+ total use-cases organized in 7 work packages (business areas), as follows:
    1. P&C Underwriting – Corporate (5 use-cases)
    2. CX Call Centre (5 use-cases)
    3. Group Life & Health (3 use-cases)
    4. Life & Health Claims (5 use-cases)
    5. Life & Health Underwriting (3 use-cases)
    6. P&C Motor Claims (1 use-cases)
    7. Technology Function (2 use-cases)

This model ensures speed, interoperability, governance, reuse, and cost efficiency, giving Ethniki the freedom to bring in the best vendor for each use case while at the same time allows candidate Vendors to bid for those case work packages, they are most experienced with.

While individual use-case work packages may be implemented to a degree in parallel to the foundational work of LOT A, they may not go into production unless they have been onboarded to the control plane.

Note on agentic use cases: Six of the twenty-four use cases in this RFP (PCM-1, CX-4, LHC-2, CPCU-4, TECH-1, TECH-2) require multi-step autonomous tool-calling workflows that go beyond standard RAG or LLM completion patterns. Where the awarded LOT A Vendor proposes the optional Agent Orchestration and Execution Layer (Component H, Section 6.3), LOT B Vendors delivering these use cases may consume Component H capabilities via the LOT A platform. Where Component H is not proposed or is not yet available, LOT B Vendors delivering these use cases must describe their intended agent orchestration approach in their technical proposal and ensure it does not introduce ungoverned shadow AI infrastructure outside the Control Plane boundary.

7.1 Unified Vendor Obligations for LOT Bx – Use-Case Delivery

The obligations defined in this section apply uniformly and in full to all LOT Bx Vendors, covering all use-case work packages across business and technology domains (LOTs B1 through B7). These obligations represent the mandatory baseline for architectural alignment, security, governance, auditability, model usage, compliance, deployment, and operational readiness.

Every LOT B Vendor must design, implement, integrate, test, deploy, and operate its solution strictly within the confines of the Enterprise AI/RAG Control Plane (LOT A) and in full alignment with the requirements and constraints set out in Sections 6, 8, 9, 10, and Annexes A, B, F, H and K. See also Section 6.1.

No LOT-specific section may supersede, contradict, weaken, or reinterpret the obligations defined herein. Any obligations listed in later sections (LOTs B1–B7) are additive only, applying solely to functional or domain-specific requirements unique to that LOT. All governance, architecture, compliance, integration, and operational requirements remain centrally defined and governed by this Section.

7.1.1 Mandatory Consumption of the Enterprise AI/RAG Control Plane

All Vendors shall:

- **Consume exclusively the LOT A Enterprise AI/RAG Control Plane** for retrieval, embeddings, model access, prompts, logging, explainability, and observability. While LOT B Vendors must consume the AI/RAG Control Plane exclusively for semantic AI functions (embeddings, retrieval, LLM inference, prompt orchestration, grounding), they would be allowed to use local deterministic or rule-based text-processing utilities (e.g., transliteration, OCR post-processing, token normalization), provided these perform no generative or learned inference. LOT B Vendors proposing a Commercially Available AI Platform (CAAP) as their primary solution must additionally comply with Section 7.9 (Certified External Platform Pathway) and obtain CEP status from the LOT A Vendor prior to promotion to UAT.

- **Not deploy or rely upon any standalone AI components**, including but not limited to private vector stores, ad-hoc embeddings, ungoverned LLM runtimes, custom retrieval layers, or shadow databases. This obligation does not apply to components of a Certified External Platform (CEP) operating under the governance bridge defined in Section 7.9, provided full CEP certification obligations are satisfied prior to UAT promotion.
- **Use only the approved Control Plane APIs** described in Annex B (Retrieval, Embedding, LLM Gateway, Prompt Governance, Audit & Explainability, Observability).
- **Respect all prohibited architectural patterns** (Annex H.9), including bypass of the gateway, unlogged inference, local caches, or any external LLM connections.

7.1.2 Functional & Technical Design Obligations

All Vendors shall deliver complete and traceable design artefacts, including:

- **Functional Specifications (FS)** covering workflows, business rules, HITL (Human-in-the-Loop) logic, exception handling and interaction patterns.
- **Technical Specifications (TS)** covering API integration, grounding strategies, retrieval behavior, prompt usage, event logging, trace IDs, and environment configuration.
- **Full alignment with Annex K – Use-Case Delivery Playbook**, including discovery artefacts, UAT criteria, test evidence, and operational runbooks.

7.1.3 Governance, Traceability & Explainability

All Vendors must ensure that every AI-assisted transaction is fully auditable and explainable. Specifically, Vendors shall:

- **Propagate all Control Plane trace IDs** and ensure that retrieval → prompt → model output → HITL decisions are reconstructible end-to-end.
- **Enable mandatory logging of prompts, retrieved chunks, model calls, safety signals, overrides, and exceptions**, consistent with Annex A and Annex H.
- **Use version-controlled prompts only** and comply with prompt-governance approval workflows.
- Ensure outputs remain **traceable to underlying source documents**, with citations and grounding evidence.

7.1.4 Security, GDPR, EU AI Act & DORA Compliance

All Vendors shall:

- Act as **GDPR Data Processors** and comply with all TOMs, minimization, retention, deletion, and data-flow obligations defined in Section 9 and Annex F1.
- Support **DPIAs, FRIAs** and AI-risk assessments for high-risk AI systems.
- Ensure full compliance with the **EU AI Act**, including dataset documentation, bias evaluation, human-oversight controls, and post-market monitoring (Annex F2).

- Implement DORA-aligned **ICT risk controls**, resilience practices, dependency statements, and operational continuity requirements (Annex F3).
- **Not export or process Ethniki data** using public AI tools, unmanaged clouds, or any non-approved environments.
- Follow all cryptographic, IAM/SSO, SIEM, network, and zero-trust requirements as defined in Section 8.

7.1.5 Deployment, CI/CD & Environment Obligations

All Vendors shall:

- **Deploy exclusively through Ethniki's CI/CD pipelines**, including mandatory security scanning, dependency checks, and environment segregation.
- Produce deployment artefacts compatible with the Control Plane environment model (DEV/TEST/PRE-PROD/PROD).
- Follow enterprise change-control processes, rollback procedures, and operational readiness steps required by Annex K.

7.1.6 Testing, Performance & Drift Management

All Vendors shall provide:

- **Test plans and evidence** across functional, integration, performance, security, explainability and bias dimensions.
- Accuracy, quality and reliability metrics specific to the use case (retrieval grounding, classification accuracy, extraction precision, automation success rates).
- **Fallback and drift-handling mechanisms**, ensuring stable behavior in cases of degraded retrieval quality, missing data, or model performance variations.
- Telemetry integration with Control Plane observability dashboards.

7.1.7 Documentation, Training & Operational Readiness

All Vendors must deliver:

- End-to-end user documentation, SOPs, troubleshooting guides, and support flows.
- Training materials for business users and operators.
- Hand-over packages meet the acceptance criteria of Annex K, Phase 6.

7.1.8 Prohibition of Ungoverned Vendor-Side Infrastructure

Vendors may **not** introduce the following, except where operating under a valid Certified External Platform (CEP) designation as defined in Section 7.9:

- external AI services or LLM endpoints
- parallel retrieval, embedding or prompt-governance layers
- shadow storage, caching, or unapproved data flows

- on-prem or cloud infrastructure outside Ethniki-approved environments
- All compute, storage, and AI operations must occur strictly within the LOT A platform *boundary* except for components of a Certified External Platform (CEP) with active CEP status under Section 7.9, which operates on vendor-hosted infrastructure subject to the data residency, audit, and DORA obligations in Section 7.9.2.

7.2 LOT B1 – P&C Corporate Underwriting – Use Case Requirements

This LOT includes the following use cases (see 13.3 Annex C – Use-Case Definitions):

- **CPCU-1** Automated Geocoding & GIS Integration
- **CPCU-2** Intelligent Mailroom Automation
- **CPCU-3** Information Aggregation & Pricing Automation
- **CPCU-4** Public Tender Intelligence & Processing
- **CPCU-5** Contract Generation & Automation

7.3 LOT B2 – CX Contact Centre – Use Case Requirements

This LOT includes the following use cases (see 13.3 Annex C – Use-Case Definitions):

- **CX-1** Transcription & CRM Population
- **CX-2** Real-Time Operator Assistant
- **CX-3** Transaction Automation
- **CX-4** Assessor Appointment Scheduling
- **CX-5** Form & Email Categorization

7.4 LOT B3 – Group Life & Health – Use Case Requirements

This LOT includes the following use cases (see 13.3 Annex C – Use-Case Definitions)::

- **GLH-1** Automated Offer Generation
- **GLH-2** Contract Generation Automation
- **GLH-3** System Integration & Member Onboarding

7.5 LOT B4 – Life & Health Claims – Use Case Requirements

This LOT includes the following use cases (see 13.3 Annex C – Use-Case Definitions):

- **LHC-1** Email Automation
- **LHC-2** Hospitalization Claims Routing
- **LHC-3** Invoice Data Extraction
- **LHC-4** Fraud Detection
- **LHC-5** Support Query Automation

7.6 LOT B5 – Life & Health Underwriting – Use Case Requirements

This LOT includes the following use cases (see 13.3 Annex C – Use-Case Definitions):

- **LHU-1** Client Entity Matching
- **LHU-2** Medical Document Summaries
- **LHU-3** Addendum Automation

7.7 LOT B6 – P&C Motor Claims – Use Case Requirements

This LOT includes the following use cases (see 13.3 Annex C – Use-Case Definitions):

- **PCM-1** Automated FNOL

7.8 LOT B7 – Technology – Use Case Requirements

This LOT includes the following use cases (see 13.3 Annex C – Use-Case Definitions):

- **TECH-1** IT / Service Desk Front-line and Back-office fulfilment operations
- **TECH-2** IT / Software Development / DevOps / SecDevOps

Vendor obligations for LOT B7:

- Vendors delivering LOT B7 use cases shall comply fully with all obligations defined in Section 7.1 — Unified Vendor Obligations for LOT Bx — including all architecture, governance, security, privacy, EU AI Act, DORA and CI/CD requirements. The obligations in Section 7.1 apply full and without exception.
- **For TECH-1 – AI-Assisted Service Desk**
    - Ensure native or seamless integration with ServiceNow ITSM for ticket creation, enrichment, classification, knowledge retrieval and automation.
    - Ensure that routing, assignment, enrichment and AI-suggested actions remain transparent and auditable and follow Service Desk operational policies.
    - Provide training and operational documentation for front-line agents and back-office contributors.
- **For TECH-2 – AI-Assisted Software Development / DevOps**
    - Integrate safely with Ethniki's software development toolchain, including IDEs, CI/CD, code repositories, and pipeline security tools.
    - Enforce secure-by-design practices including SAST/DAST, dependency scanning, secret management, and protected code-generation workflows.
    - Ensure that no AI features bypass version control, code review or mandatory pipeline quality gates.

7.9 Certified External Platform (CEP) Pathway

The obligations in Section 7.1 apply in full to all LOT B Vendors. Where a LOT B Vendor proposes a Commercially Available AI Platform (CAAP) as the primary delivery solution for a use-case work package, the Vendor may seek Certified External Platform (CEP) status for that platform. The CEP pathway does not reduce any obligation on the LOT B Vendor; it provides a structured governance bridge enabling commercially mature platforms to integrate with the LOT A Control Plane under full audit authority, rather than being categorically excluded. The LOT B Vendor remains fully responsible for all use-case delivery outcomes regardless of CEP status.

7.9.1 Eligibility Criteria for CEP Designation

A commercially available platform is eligible for CEP consideration if it satisfies all of the following criteria at the time of bid submission:

- The platform is in production operation at a minimum of two regulated European insurance organisations, with verifiable written references available upon request during the evaluation process.
- The platform's primary AI functions are directly relevant to the LOT B use-case work package for which it is being proposed.
- The platform vendor is capable of entering a GDPR-compliant Data Processing Agreement with Ethniki Asfalistiki designating the platform as a data processor under Article 28 GDPR.
- The platform vendor accepts the technical integration requirements defined in Section 7.9.2 as a binding condition of CEP status, and commits to completing those integrations prior to UAT promotion.

7.9.2 Mandatory CEP Integration Obligations

To obtain and maintain CEP status, a platform vendor must implement all seven of the following obligations within the LOT A technical environment. These obligations are verified by the LOT A Vendor as part of the technical certification process and are prerequisites for UAT promotion:

- **Audit Log Export.** The platform shall export a complete, real-time structured log of all AI inference events (LLM calls, document processing, retrieval queries, agent actions, HITL events) to the LOT A audit ingestion API. Log format shall comply with the LOT A audit schema specification. Continuous monitoring post-deployment is required.
- **LLM Gateway Compliance or Bridge.** Any generative AI functions within the platform shall route LLM calls through the LOT A LLM Gateway, or demonstrate equivalent governance via a LOT A-approved bridging integration. Direct calls to model endpoints that bypass LOT A's Gateway are prohibited unless a technical exemption is granted in writing by Ethniki on the recommendation of the LOT A Vendor.
- **Data Residency Attestation.** The platform vendor shall provide written attestation that all data processing occurs within EU jurisdiction, shall accept Ethniki's right to audit data residency controls at any time during the contract period, and shall notify Ethniki within five business days of any change to data residency arrangements.

- **DORA Dependency Mapping.** The platform shall be included in Ethniki's DORA ICT third-party risk register. The platform vendor shall provide an ICT concentration risk disclosure, a contractually binding business continuity and exit plan, and shall participate in periodic resilience testing as required by Annex F3.
- **EU AI Act Risk Classification.** The platform vendor shall provide a completed AI system risk classification under EU AI Act requirements. For use cases assessed as high-risk, the platform vendor shall support Ethniki in completing the required Fundamental Rights Impact Assessment (FRIA) and shall make available all documentation required under Annex F2.
- **Observability Integration.** The platform shall expose platform-level telemetry (availability, latency, error rates, throughput) to the LOT A Dynatrace observability instance as specified in Section 6.2(F), enabling unified operational monitoring alongside native LOT A components.
- **Incident and Change Notification.** The platform vendor shall provide advance notification of any significant platform changes (model updates, architecture changes, security patches affecting AI inference) to the LOT A Vendor, with a minimum lead time of ten business days except for emergency security patches. All changes shall be recorded in Ethniki's change register and subject to LOT A certification review before promotion to PROD.

7.9.3 LOT A Vendor Responsibilities for CEP Governance

The LOT A Vendor shall manage the CEP programme on behalf of Ethniki. LOT A responsibilities include: conducting the technical certification assessment for each candidate platform within thirty business days of a complete application; maintaining the CEP register listing all platforms that have achieved, failed, or had CEP status revoked; performing quarterly compliance reviews of all active CEP platforms including audit log spot checks and observability data review; exercising LOT A blocking rights in respect of any CEP platform that fails a compliance review; and reporting CEP compliance status to Ethniki's programme governance on a quarterly basis. Certification timelines shall not create grounds for LOT B schedule relief unless delay is demonstrably attributable to LOT A.

7.9.4 CEP Evaluation and Scoring

- LOT B proposals incorporating a CEP platform shall be evaluated under the standard Section 11.2 criteria. A CEP platform's proven production deployment record at regulated European insurers shall be treated as equivalent evidence to a custom build's technical architecture score under the "Functional Accuracy & Completeness" criterion, provided all CEP certification obligations in Section 7.9.2 are accepted in full at the time of bid submission.
- LOT B Vendors shall disclose at bid submission whether they intend to seek CEP status for any proposed platform, including the platform name, platform vendor, and a preliminary CEP compliance gap assessment against Section 7.9.2. Failure to disclose a CAAP at bid stage and subsequent discovery of undisclosed platform AI infrastructure during the LOT A certification process may constitute a material compliance failure.

- CEP designation does not reduce the LOT B Vendor's contractual accountability for use-case delivery outcomes. The LOT B Vendor remains fully responsible for the performance, compliance, and operational behaviour of any CEP platform within their work package, including platform failures, data quality issues, and regulatory compliance gaps.

8 Information Security & DORA Requirements

Ethniki Insurance requires all Vendors to comply with enterprise-grade Information Security, Cybersecurity and ICT Risk Management obligations aligned with ISO 27001, ISO 27002, NIST CSF, GDPR, PCI-DSS and DORA.

The Vendor shall, at a minimum:

- Implement a zero-trust security model with strict least-privilege and role-based access control (RBAC) across all components of the solution, including administrative tooling.
- Ensure strong authentication (including MFA) integrated with Ethniki's enterprise Identity and Access Management (IAM) platform for all user, service and API access.
- Encrypt all data in transit using TLS 1.2 or higher and all data at rest using AES-256 or equivalent quantum resistant cryptographic systems, including vector stores, logs, backups and model artefacts.
- Apply secure configuration baselines and hardening standards to all infrastructure components, runtimes and managed services, whether hosted on-premise or in the cloud.
- Will NOT store, process, or transmit Ethniki Insurance's (EI) PII in any public AI tools or LLMs, outside of EI's corporate on-prem and cloud environments.
- Consistent performance (i.e., meet predefined SLAs, KPIs) for all AI deliverables and proper design and architecture to ensure they will scale well.

LOT A – Enterprise AI/RAG Control Plane shall additionally:

- Provide centralized logging and full audit trails for all relevant security events, including authentication, authorization, configuration and access role/rights changes, data access, prompts, retrieved context, model invocations and generated outputs.
- Integrate with Ethniki's Security Information and Event Management (SIEM) solution in near real-time and support alerting configurable thresholds and detection rules.
- Enforce technical segregation between environments (development, test, pre-production, production) and between tenants/use-cases where applicable.

All Vendors (LOT A and LOT Bx) must implement and evidence a Secure Software Development Life Cycle (SSDLC), including:

- Static and dynamic application security testing (SAST/DAST) and dependency scanning as part of CI/CD pipelines.
- Regular vulnerability management and patching processes with agreed SLAs for remediation of critical and high-severity findings.

- Strict control of secrets and keys (e.g. use of vaulting solutions) and prohibition of secrets in source code.

In line with DORA obligations, the Vendor shall:

- Implement and maintain documented ICT risk-management practices covering identification, assessment, mitigation and monitoring of ICT risk relevant to the provided services.
- Deliver design and architecture that will support our Business Continuity (BCP) and Disaster Recovery (DRP) Plans, for all critical components supporting the services, and meet Recovery Time Objective (RTO) ≤ 4 hours and Recovery Point Objective (RPO) ≤ 0 minutes, unless otherwise agreed.
- Support periodic resilience testing (including DR tests, fail-over tests and other scenario-based exercises) and provide test reports to Ethniki upon request.
- Notify Ethniki of any material ICT-related incident affecting the services without undue delay and within the maximum timelines contractually agreed, and support Ethniki in fulfilling any regulatory reporting obligations.

No shadow IT or unapproved infrastructure is permitted. Vendors shall not store or process Ethniki data on local devices or in non-approved cloud services. All compute and storage resources used for delivering the Services must be declared, approved and subject to the above security and DORA requirements. Note, CEP platforms satisfy the "declared and approved" requirement through the CEP certification process under Section 7.9; their infrastructure remains subject to the data residency attestation and DORA dependency mapping obligations in Section 7.9.2.

9 Data Privacy & GDPR Requirements

All Vendors shall act as Data Processors under GDPR, with Ethniki Insurance acting as Data Controller. All personal data processing carried out in the context of the Services must be strictly necessary, proportionate and aligned with the purposes defined by Ethniki.

The Vendor shall ensure full alignment with the core GDPR principles (lawfulness, fairness, transparency, purpose limitation, data minimization, accuracy, storage limitation, integrity, confidentiality, availability and accountability) and shall implement appropriate technical and organizational measures (TOMs) to ensure the security of the data, including the personal data, processed on behalf of Ethniki in accordance with Articles 5, 6, 9, 28 and 32 of GDPR. These measures must consider the latest technological advancements, the costs of implementation, the nature, scope, context and purpose of processing, as well as the risk associated with the processing activities.

In particular, the Vendor shall:

- Process personal data solely on documented instructions from Ethniki.

- Ensure that only the minimum personal data required for the specific processing purpose are ingested into, or exposed via, the solution, and that privacy-by-design and privacy-by-default principles are applied throughout.
- Ensure that its employees, as well as all other persons engaged in the project and authorized to access personal data, have committed themselves to confidentiality and to processing the data exclusively in accordance with the instructions provided by Ethniki.
- To ensure appropriate technical measures, as outlined above, the Vendor shall, as appropriate, ensure the following: a) continuous confidentiality, integrity, availability and resilience of processing systems and services; b) the ability to promptly restore the availability and access to personal data in a timely manner in the event of a physical or technical incident; c) a process for regularly testing, assessing and evaluating the effectiveness of technical and organizational measures in place to ensure the security of data processing.
- Data localization and legal sovereignty: The Vendor shall guarantee that Ethniki's Data will remain strictly under national or European jurisdiction, with no possibility of foreign interference.
- There must be total transparency regarding storage locations and any subcontracting chain. The Vendor shall ensure that the Data may be transferred only to its subcontractors located in member countries of the European Economic Area and/or third countries recognized by the European Commission as providing an adequate level of protection, pursuant to an adequacy decision.
- Implement TOMs covering access control, encryption, pseudonymization / anonymization (where applicable), data integrity, backup and recovery, secure deletion, logging and monitoring as described in Annex F1 – GDPR Compliance Pack.
- Maintain and provide, upon request, records of processing activities and support Ethniki in the performance of Data Protection Impact Assessments (DPIAs) where processing - in particular using new technologies, and taking into account the nature, scope, context and purposes of the processing - is likely to result in a high risk to the rights and freedoms of natural persons.
- Ensure that any sub-processor is engaged only with prior written authorisation from Ethniki and is bound by data protection obligations no less protective than those imposed on the Vendor.
- Any kind of transferring of personal data outside the EEA without Ethniki’s prior approval and written consent and without implementing an appropriate transfer mechanism (e.g. Standard Contractual Clauses) is strictly prohibited.

Vendor solutions must support, through configuration and/or APIs, the execution of data subject rights (access, rectification, restriction, portability, deletion/erasure) without compromising auditability or platform integrity. LOT A shall enable centralised tooling to support such requests where feasible.

Personal data retention periods must be configurable and controlled by Ethniki. Vendor shall ensure, and be able to demonstrate at any time, that personal data are deleted or irreversibly

anonymised at the end of the applicable retention period, or upon Ethniki's instruction, including backups and derived artefacts within agreed timelines.

Vendors are not permitted to introduce standalone databases, vector stores, embedding repositories or local caches containing Ethniki data outside the approved Enterprise AI/RAG Control Plane, except where those components form part of a Certified External Platform (CEP) operating under the governance obligations of Section 7.9. All data flows and storage locations, including those within CEP-designated platforms, must be documented and visible to Ethniki, as further specified in Annex F1 – GDPR Compliance Pack.

Ethniki reserves the right to conduct any checks it considers useful to ascertain the Vendor's fulfilment of his obligations, particularly through an audit. The Vendor undertakes the responsibility to reply without unnecessary delay to Ethniki's requests for audits and support the process throughout.

10 EU AI Act Compliance Requirements

Many of the AI/IDP solutions procured under LOT A and LOT Bx are expected to fall within the scope of, or be designed to meet the requirements of, High-Risk AI systems under the EU AI Act. Vendors must therefore adopt an EU AI Act-aligned approach to the design, development, deployment and operation of their solutions.

The Vendor will implement and maintain an AI risk-management system covering the entire lifecycle of the AI components, including identification and analysis of risks to fundamental rights, health and safety, as well as risks associated with data quality, bias, robustness and cybersecurity. Mitigation measures and residual risks shall be documented and kept up to date.

The Vendor will ensure that training, validation and testing datasets used for AI models are relevant, representative, free of known errors to the extent reasonably possible and appropriately documented. Data governance practices (e.g. source documentation, labelling, pre-processing, augmentation, bias assessments) must be described and made available to Ethniki as part of Annex F2 – EU AI Act Compliance Pack.

The Vendor shall provide a Fundamental Rights Impact Assessment (FRIA) for any use case that directly or indirectly affects individuals' rights, freedoms, or entitlements, in accordance with EU AI Act Articles 9–29. The FRIA shall include assessment of potential harm, mitigations, and validation of fundamental rights protections.

LOT A – Enterprise AI/RAG Control Plane must provide technical capabilities to support EU AI Act compliance, including at a minimum:

- Comprehensive logging and traceability of inputs, retrieved context, model parameters and outputs for all AI-enabled interactions.

- Versioning and lifecycle management of models, prompts, embeddings and configuration, with clear rollback and change-control mechanisms.
- Explainability and transparency tooling appropriate to the risk level and use case, enabling Ethniki to understand the main drivers behind AI-generated outputs and to provide meaningful explanations to affected individuals where required.
- Monitoring for model and embedding drift, and for potential bias or performance degradation, with alerting and re-training workflows.
- Safe-completion controls and guardrails to prevent obviously harmful, discriminatory or non-compliant outputs.

Each Vendor delivering LOT Bx use cases shall provide, as applicable:

- A description of the AI models used (including third-party or General-Purpose AI (GPAI) services), their intended purpose and their integration within the business processes.
- Documentation of the datasets used for training/finetuning and for validation, including data provenance, quality checks and known limitations.
- Accuracy, robustness and cybersecurity evaluations (including adversarial robustness, where relevant) and the associated monitoring plan.
- A human oversight concept describing how Ethniki personnel will remain in control of, and able to override, AI outputs in decision workflows.
- A post-market monitoring concept for the AI components, including incident logging and feedback loops from users.

Any use of third-party GPAI services or external LLMs must be explicitly declared. Such services may only be accessed via the LOT A Control Plane or, where the LOT B Vendor is deploying a Certified External Platform (CEP) under Section 7.9, via the CEP governance bridge approved by the LOT A Vendor. Direct connections from use-case solutions to external AI services that are neither routed through the LOT A Control Plane nor covered by an active CEP designation are not permitted. Vendors shall cooperate with Ethniki in preparing and maintaining any documentation required for conformity assessments or regulatory supervision, as further detailed in Annex F2 – EU AI Act Compliance Pack.

11 Technical Evaluation Criteria

Evaluation of submitted solutions will be based on a set of technical and functional criteria designed to ensure alignment with organizational needs and regulatory requirements. These criteria are segmented into LOT A and LOT Bx categories, each with specific focus areas and weighted percentages to guide the assessment process. By adhering to these criteria, vendors can demonstrate the quality, reliability, and compliance of their solution offerings within the defined framework.

11.1 Evaluation Criteria for LOT A (Control Plane Vendor)

Criterion	Description	Reference Section(s)
Technical Solution – Control Plane Architecture	Assessment of platform design including retrieval, vector store, gateway, ingestion, observability.	Sections 6, Annex A, Annex H
Security, Compliance & Governance	Evaluation of GDPR, EU AI Act, DORA alignment, audit readiness & security posture.	Sections 8–10, Annex F
Operational Readiness & Scalability	Resilience, DR/BCP, SLA commitments, monitoring and AI-Ops maturity.	Section 6.1(F–G), Annex H
Vendor Innovation & Delivery Model	Capability to deliver enterprise-grade AI platforms at scale; quality of delivery plan. Proposals including the optional Agent Orchestration and Execution Layer (Component H, Section 6.3) will be evaluated on agentic governance maturity, roadmap credibility, and ability to support the six identified agentic use cases under this criterion.	Various, Annex K

11.2 Evaluation Criteria for LOT Bx (Use-Case Vendors)

Criterion	Description	Reference Section(s)
Functional Accuracy & Completeness	Quality of use-case design, workflow alignment, handling of business rules and exceptions.	Section 7, Annex L
Integration with LOT A	Correct usage of APIs, retrieval, prompts, explainability, HITL workflows.	Section 6, Annex B, Annex H
Security, Compliance & Data Handling	GDPR-by-design processing, EU AI Act requirements, DORA operational dependency mapping.	Sections 8–10, Annex F
Vendor Innovation & Delivery Model	Quality of delivery methodology, team composition, timeline feasibility.	Various, Annex K

Note, references to Section 6 are only integration related.

12 Legal & Contractual Requirements

The following requirements apply to all Vendors participating in this RFP and shall form part of the contractual framework governing LOT A (Enterprise AI / RAG Control Plane) and LOTs B1–B7 (AI and IDP Use-Case Delivery).

Where obligations relate to model lifecycle, retrieval, ingestion, auditability, or platform configuration, these apply to Vendors according to the responsibilities of their awarded LOT, as defined in Sections 6 and 7. LOT B Vendors do not deliver or operate platform components; they only consume approved LOT A services. Where the Vendor does not deliver a model but consumes the LOT A LLM Gateway, 'model-related obligations' apply to the workflow, prompts, retrieval configuration, and quality metrics of the LOT B deliverable.

12.1 Contracting Structure

The contractual framework shall consist of a Master Services Agreement (MSA), individual Statements of Work (SOWs) per LOT, Annex F1–F3 compliance packs, pricing schedules, and Service Level Agreements (SLAs).

Additional provisions:

- The allocation of deliverables and responsibilities between LOT A and LOTs B1–B7 shall be explicitly defined, including dependencies, integration obligations and cross-vendor interfaces.
- The contracting model shall prevent responsibility gaps in a multi-vendor environment and establish a clear escalation and issue-resolution mechanism across Vendors.
- Ethniki may, at its discretion, introduce additional SOWs (e.g., Support & Maintenance) prior to production.

12.2 Data Protection & Regulatory Compliance

Vendor shall comply with GDPR, EU AI Act, PCI-DSS, DORA and all regulatory requirements applicable to Ethniki's operations, and shall accept Annex F as binding.

Additional provisions:

- Vendors shall not utilize any unapproved external AI endpoints or telemetry services. This restriction does not apply to: (a) deterministic, non-AI utilities (e.g., OCR post-processing, text normalization, rule-based parsing) that perform no generative, learned or probabilistic inference; or (b) AI components forming part of a Certified External Platform (CEP) that holds active CEP status under Section 7.9 and satisfies all mandatory integration obligations therein, including audit log export and LLM Gateway compliance or bridge.

- Vendor shall maintain Model Cards, Data Sheets and full lifecycle documentation for all delivered models or model-based workflows delivered or configured under the awarded LOT.
- Vendor shall ensure comprehensive audit logging of LLM Gateway inference events, retrieval, access and versioning activities.

These obligations complement, and do not supersede, the technical and compliance requirements defined in Sections 8–10 and Annex F. In case of contradiction, Sections 8–10 and Annex F prevail for compliance specifics.

12.3 Intellectual Property Rights

All deliverables produced for Ethniki—shall be owned by or licensed to Ethniki for unrestricted enterprise use.

Additional provisions:

- Vendor pre-existing IP shall remain the property of the Vendor; however, Vendor shall grant Ethniki a perpetual, irrevocable, royalty-free enterprise license to use, modify, extend and integrate such IP without limitation.
- Vendor shall not use Ethniki data, prompts, logs or configurations to train, enhance or enrich products or services for third parties.
- Where Vendor proprietary components are critical to service continuity, an escrow mechanism may be required at Ethniki's discretion.

This does not apply to standalone commercial off-the-shelf software or pre-existing platform components that are licensed independently and not embedded into Deliverables.

12.4 Confidentiality

Vendor shall protect all Ethniki information and maintain strict data isolation.

Additional provisions:

- No Ethniki data, metadata or artefacts may be used for external training, evaluation, benchmarking or optimisation.
- Vendor shall not retain or export LLM telemetry or user interaction logs unless expressly authorised in writing.

12.5 Sub-contracting & Sub-processors

Vendor shall not engage sub-contractors or sub-processors without Ethniki's prior written approval and shall ensure equivalent obligations flow down to all third parties.

Vendor shall disclose all cloud regions, model hosting locations, infrastructure providers, LLM endpoints and vector database services used in service delivery.

Sub-processor requirements shall be interpreted together with Annex F1 (GDPR Compliance Pack).

12.6 Security & Audit Rights- Regulatory Cooperation

Ethniki retains full audit rights covering technical, operational and data protection compliance. Vendors shall support regulator audits and provide access required for supervisory reviews, shall support responses to supervisory authority inquiries, provide explainability documentation, and assist in regulatory model validation activities.

Vendors must also comply with the platform-specific architectural, security, and observability requirements defined in Annex H.

12.7 Service Levels & Service Credits

SLAs shall include availability, performance, accuracy and incident response requirements.

Additional provisions:

- AI-specific KPIs shall include model drift thresholds, hallucination tolerances, retrieval precision/latency metrics and IDP accuracy benchmarks (e.g., F1 scores).
- Breaches may trigger service credits, mandatory remediation actions or re-validation of the affected model or workflow.

The SLAs in this Section complement—do not override—the technical and architectural SLAs specified in Sections 6 and 7.

12.8 Change Control

All changes to datasets, parameters, models or model-based workflows delivered or configured under the awarded LOT, shall follow a formal Change Control Process, including:

- impact assessment and risk classification,
- testing, validation and sign-off,
- rollback and fallback procedures,
- regulatory impact analysis where required under the EU AI Act.

No changes may be promoted to production without written approval by Ethniki. LOT A Vendor shall conduct technical certification of LOT B changes before Ethniki grants final approval.

12.9 Exit, Transition & Portability

Vendor shall ensure an orderly and complete transition to Ethniki or a replacement Vendor. Obligations include:

- full export of all applicable model assets (for LOT A only), embeddings, vector databases, training artefacts, prompts, orchestrations and pipelines in open and interoperable formats. LOT B Vendors export only the artefacts they deliver (e.g., workflows, prompts, configurations). LOT A Vendor remains responsible for the export of platform-level assets such as vector databases, embeddings, model-routing configurations, and control-plane pipelines.
- delivery of complete technical and operational documentation,
- provision of transition assistance for a period defined by Ethniki (minimum 3–6 months),
- issuance of data deletion and sanitization certificates covering production and backup environments.

Ethniki must retain the ability to continue operating or migrating the solution independently of the Vendor.

12.10 Indemnities & Liability

Vendor shall indemnify Ethniki for:

- data protection breaches,
- IP infringement,
- regulatory fines arising from Vendor acts or omissions,
- AI-related harms, including model defects, biased outputs, hallucinations and automation errors affecting customer, underwriting or claims decisions.

Vendor may not rely on “algorithmic unpredictability” as a liability exclusion. Liability caps shall reflect the criticality of enterprise-wide AI operations.

12.11 Insurance Requirements

Vendor must have cyber, professional liability, and business continuity insurance covering AI-related errors, data poisoning, model integrity issues, and operational disruptions to core processes.

12.12 Additional Requirements

- **Algorithmic Transparency**: Vendor shall provide, upon request, model architecture documentation, training data provenance summaries, explainability artefacts, retrieval traceability outputs and evaluation results. Technical certification of LOT B deliverables shall follow the LOT A governance framework defined in Section 6.1.
- **Vendor Performance & Key Personnel**: Vendor shall maintain adequate staffing and expertise and shall not replace key personnel without Ethniki approval.
- **Conflicts of Interest**: Vendor shall disclose potential conflicts, including engagements with competing insurers. Ethniki may require implementation of ethical walls.

- **Business Continuity & Disaster Recovery**: Vendor shall maintain and test DR plans covering AI components, in alignment with DORA resilience requirements.

13 Annexes

13.1 Annex A – Control Plane Requirements (LOT A only)

This annex defines the **capabilities**, **governance standards**, **security requirements**, and **operational commitments** that the LOT A Vendor must deliver and maintain throughout the engagement. The solution is critical as it acts as the shared AI foundation for all vendors and all use cases. The capabilities described here are binding and supersede any vendor-specific interpretations unless explicitly agreed in writing. As such, the requested items combine professional services with technological elements.

13.1.1 Purpose of the Control Plane

The Control Plane shall:

- Provide a **single, governed AI foundation** for all enterprise use cases.
- Enable consistent, explainable, and compliant operation of generative AI systems.
- Serve as the mandatory integration point for all LOT Bx vendors.
- Ensure that AI services remain secure, resilient, observable, and fully auditable.
- Centralize model access, retrieval, prompt governance, performance management, and regulatory evidence creation.

The Control Plane is the **only permitted environment** for the following functions, except where those functions are performed by a Certified External Platform (CEP) operating under the governance bridge defined in Section 7.9:

- Retrieval-Augmented Generation workflows
- Vector indexing and embedding storage
- Enterprise LLM access and routing
- Governance, logging, auditability, and explainability functions
- Prompt management and version control
- Vendor onboarding and supervised AI operations

The following are all **mandatory** unless specified otherwise.

13.1.2 Retrieval & Vector Search Layer

A consistent, high-quality semantic retrieval backbone shared across all vendors and use cases.

- Dense, sparse and hybrid retrieval mechanisms with configurable scoring.
- Scalable vector indexing optimized for enterprise workloads.
- RBAC-aware retrieval ensuring segment-appropriate and role-appropriate access.

- Metadata/tag-based filtering for precision use cases.
- Full lineage from source → chunk → embedding → index → retrieval result.
- Retention, expiry and index reconstruction policies aligned with governance rules.
- Retrieval quality monitoring: precision@k, hallucination correlation, grounding
- Outcome:

13.1.3 Ingestion, Pre-Processing & Embedding Services

Clean, structured and searchable enterprise knowledge — a prerequisite for safe RAG.

- Ingestion of PDFs, emails, messages, DOCX, XLSX, images, and scanned documents.
- OCR/ICR, layout recognition, table extraction and structural parsing.
- Semantic chunking strategies configurable per use case.
- Metadata enrichment for document provenance, confidentiality and lifecycle.
- Embedding generation using enterprise-approved embedding models or model or configured under the awarded LOT.
- Re-embedding workflows triggered by source updates.
- Validation checks ensuring embedding consistency and vector health.

13.1.4 LLM Gateway & Model Routing

The platform acts as the single controlled gateway to all LLMs and generative models to ensure safe, governed, consistent use of generative AI across Ethniki Asfalistiki.

- Routing to multiple approved LLMs (commercial, OSS, internal).
- Policy-based routing (cost, latency, sensitivity, safety class, throughput).
- Enforced grounding and hallucination-mitigation logic.
- Parameter controls (temperature, max tokens, top-p, sampling constraints).
- Redaction and input-sanitization, where required.
- Logging of every model invocation with full context.
- Controlled access to third-party or remote LLM endpoints; no direct vendor calls permitted.

13.1.5 Prompt Management & Governance Layer

The platform needs to deliver robust enterprise prompt governance, effectively treating prompts as governed enterprise assets, not ad-hoc artefacts.

- Version-controlled prompt repository with approval workflows.
- Parameterized prompts supporting use-case-specific variables.
- Guardrail enforcement (blocked terms, response shaping, safety filters).
- Prompt lineage tracking linking prompts to outputs, datasets, and policies.
- A/B testing and evaluation workflows.
- Prompt-level access controls and audit trails.

13.1.6 Governance, Auditability & Explainability

To comply with GDPR, DORA and EU AI Act, the Control Plane needs to behave and enforce an audit ready posture maintaining regulator-ready transparency at all levels and for all onboarded use cases.

- Full logs of prompts, retrieved chunks, model decisions, outputs and overrides.
- Time-sequenced audit trails for investigations and regulatory requests.
- Built-in explainability functions enabling interpretation of generated responses.
- Retrieval citations showing source documents and relevance weightings.
- Supplier-agnostic AI Risk Pack generation for LOT Bx use cases.
- Support for DPIA, model risk assessment and EU AI Act technical documentation.

13.1.7 Monitoring, Observability & AI-Ops

The platform must be in a position to enable and enforce continuous monitoring ensuring a measurable, stable and predictable AI operation for our enterprise.

- Retrieval metrics: precision, recall, grounding quality, fallback rate.
- Model metrics: latency, error rates, token consumption, throughput.
- Drift detection: embedding drift, retrieval drift, model drift and prompt drift.
- Usage dashboards aggregated by vendor, use case and department.
- Costs & benefits aggregated by vendor, use case and department.
- SIEM integration for real-time anomaly and threat detection.
- Alerting thresholds defined per use case and risk level.
- Observability aligned to Ethniki Asfalistiki observability stack.

13.1.8 Environment Management, CI/CD & Operational Readiness

The platform needs to provide enterprise-level resilience, able to scale for the use-cases of this RFP and many more to come.

- Segregated DEV / TEST / PRE-PROD / PROD environments.
- Automated deployment pipelines, including vendor onboarding workflows.
- Rollback and version pinning for models, prompts, embeddings and indexes.
- Backup/restore mechanisms aligned with DORA (RTO ≤ 4 hours, RPO ≤ 0 minutes).
- Disaster Recovery readiness across all critical components.
- Performance Baselining and capacity auto-scaling.

13.1.9 Vendor Onboarding & Multi-Vendor Integration Requirements

All LOT B vendors plug into the same governed backbone. The Vendor solution needs to effectively and efficiently include services (QA) and technological elements to support the development, roll-out and operations stipulated in this RFP.

- Provide a standard onboarding playbook for LOT B vendors or vendors of future solutions.
- Expose well-documented, stable APIs with schemas and examples.
- Support per-vendor workspaces or isolation boundaries when required.
- Enforce consistent security, logging, embedding, and retrieval behaviors across vendors and use-cases.
- Guarantee that no LOT B vendor can introduce shadow LLMs or ungoverned retrieval layers. Where a LOT B vendor proposes a Certified External Platform (CEP) under Section 7.9, the LOT A Vendor is responsible for conducting the CEP certification assessment, maintaining the CEP register, and enforcing all CEP integration obligations as part of the governance programme.
- Provide cross-vendor / cross-use-case debugging and inspection tools.

13.1.10 Operational & Service Management Requirements

All LOT B vendors plug into the same governed backbone making it a critical component of Ethniki's AI/IDP & Automations Delivery Framework; thus, the Vendor solution needs to guarantee enterprise-grade operational assurance. Support services will be provided by an additional support and maintenance engagement to be confirmed in due course and before elements of the framework go into production.

- 24/7 support for critical incidents.
- Defined SLAs for uptime, latency, retrieval and API availability.
- Incident classification, escalation and RCA processes.
- Monthly operational reporting and quarterly performance review.
- Change control aligned to ITIL and regulatory requirements.
- Vendor resource continuity and competency obligations.

13.1.11 Data & Content Lifecycle Governance

To ensure controlled, compliant, lifecycle-aware AI data operations, the following are mandatory:

- Retention and expiry policies for embeddings, logs, and retrieved data.
- Automatic handling of deleted/updated source documents.
- GDPR-aligned deletion, correction and access request support.
- Version-level lifecycle control for chunks, embeddings, prompts and models.
- Data cataloguing and lineage visibility at each processing step.

13.1.12 Mandatory Deliverables for LOT A Vendor

Vendors must deliver and maintain tangible artefacts enabling governance, auditability and cross-vendor integration:

- High-level & low-level platform architecture diagrams.
- Vector index design, embedding strategy, ingestion pipelines.
- Prompt governance framework and prompt library baseline.
- API documentation, SDKs, integration examples.
- Compliance artefacts (GDPR, EU AI Act, DORA).
- Security architecture, pen test reports, hardening standards.
- Operational runbooks, DR plans, on-call procedures.
- Monitoring dashboards, drift reports, capacity plans.
- Onboarding kits for LOT B vendors.

13.1.13 Optional Value-Add Platform Capabilities

The following are not mandatory but may positively differentiate proposals:

- Retrieval acceleration (FAISS/HNSW/ColBERT).
- Multilingual embedding support.
- Document clustering, summarisation, similarity analytics.
- Knowledge quality scoring and enrichment.
- Controlled experimentation workspace for new prompts/flows.
- Automated dataset profiling and content health assessment.
- Semantic monitoring of drift or knowledge gaps.

13.2 Annex B – API Specifications

13.2.1 Purpose

Annex B may not contain concrete specifications since **LOT A will design and deliver the actual platform and LOT B vendors will later integrate with it**. Hence, we are specifying mandatory API behaviors, shapes, boundaries, governance rules, and performance expectations that **LOT A MUST deliver and LOT B MUST consume**. This is the standard approach used in platform RFPs. In addition, in the context of this RFP, this approach enables Ethniki to define API categories, capabilities, interaction models, security rules, versioning policies, and lifecycle requirements for integration between LOT A and LOT B solutions. Specifically, it:

- Enables pricing.
- Avoids over-specification.
- Protects architecture flexibility.
- Creates enforceable obligations.
- Ensures that LOT B vendors know what integration looks like.
- Ensures that LOT A vendor knows what they are committing to deliver.

13.2.2 Retrieval API (LOT A Must Deliver) Capabilities & LOT B Obligations

Capabilities:

- Inputs: query text, metadata filters, user identity, grounding settings.
- Outputs: ranked passages, metadata, confidence scores, trace IDs.
- Behaviors: hybrid retrieval (dense + sparse), RBAC-based filtering, full logging.

LOT B obligations:

- Must not implement custom retrieval or embeddings.
- Must propagate trace IDs.

13.2.3 Ingestion & Embedding API Capabilities & Behaviors & LOT B Obligations

Capabilities:

- Multi-format ingestion (PDF, DOCX, XLSX, images).
- OCR, metadata extraction, semantic chunking.
- Automatic embedding generation.

Behaviors:

- Async ingestion; job polling for completion.

LOT B obligations:

- Cannot modify index or direct-write embeddings.

13.2.4 LLM Gateway API Capabilities & Outputs & LOT B Obligations

Capabilities:

- Model routing by cost, latency, sensitivity.
- Mandatory grounding.
- Safety filtering and redaction.

Outputs:

- Structured response, safety flags, confidence.

LOT B obligations:

- No direct API calls to external LLMs.

13.2.5 Prompt Management API Capabilities & LOT B Obligations

Capabilities:

- Retrieve approved prompts.
- Parameter binding.
- Version tracking.

LOT B obligations:

- No unapproved or external prompts.
- Must use versioned prompts only.

13.2.6 Audit & Explainability API Capabilities & LOT B Obligations

Capabilities:

- Retrieval of logs, citations, embeddings used, reasoning traces.
- FRIA-relevant traceability.

LOT B obligations:

- Must expose HITL and exception handling to audit pipeline.

13.2.7 Observability & Monitoring API Capabilities & LOT A/B Obligations

Capabilities:

- Latency, drift, throughput, error metrics.
- Model usage patterns.

LOT A obligations:

- Provide dashboards and export endpoints.

LOT B obligations:

- Supply use case telemetry for monitoring.

13.2.8 API Security Requirements

- OAuth2/OIDC enforced.
- All calls routed through enterprise API Gateway.
- No unencrypted calls.
- Strict prohibition of shadow LLMs, local vector DBs, or unapproved retrieval logic, except within a Certified External Platform (CEP) holding active CEP status under Section 7.9 and operating under the full set of CEP integration obligations.

13.2.9 Performance & Reliability

- Retrieval latency < 500 ms.
- LLM gateway < 2 seconds.
- Job ingestion SLA 95% < 5 min.
- Availability 99.5%.
- DR: RTO ≤ 4h, RPO ≤ 0 min.

13.2.10 API Versioning & Lifecycle

- Semantic versioning required.
- Minimum 6-month deprecation window.
- LOT A provides migration guides.

13.2.11 Sandbox & Testing Toolkit (LOT A Must Deliver)

- Mock APIs.
- Payload examples.
- Postman collections.
- Error simulation tools.

13.2.12 Schema Finalization

- Final schemas will be developed by LOT A post-award. Vendors shall budget for REST/JSON integration with OpenAPI definitions provided during onboarding.

13.2.13 Optional Agent APIs (Required if Component H is Proposed)

Where the LOT A Vendor proposes the optional Agent Orchestration and Execution Layer (Component H, Section 6.3), the following three APIs must also be delivered. They are not required where Component H is not proposed. LOT B Vendors consuming Component H shall use these APIs for all agent lifecycle management, tool access, and audit retrieval operations.

1. **Agent Orchestration API.** Lifecycle management for agent deployments and execution sessions. Key operations: RegisterAgent, StartSession, GetSessionState, SubmitHITLResponse, TerminateSession, GetExecutionTrace. Authentication: Agent Identity plus scoped OAuth2 token. LOT B obligation: all agentic workflow sessions must be initiated and managed via this API; no direct instantiation of agent runtimes outside the LOT A platform boundary.
2. **Tool Registry API.** Registration and discovery of approved tools and external integrations available for agent use. Key operations: RegisterTool, GetToolManifest, ListApprovedTools, RevokeToolAccess, GetToolExecutionLog. Authentication: LOT A admin role for registration; LOT B service principal for discovery and invocation. LOT B obligation: agents may only invoke tools listed in the registry with explicit caller permission; unapproved tool calls shall be blocked at the sandbox boundary.
3. **Agent Audit API.** Retrieval of agent execution records for governance, compliance, and regulatory reporting. Key operations: GetAgentTrace, ListAgentSessions, ExportAuditLog, GetHITLHistory. Authentication: read-only access restricted to compliance team, internal audit, and Ethniki-authorised regulators. LOT B obligation: LOT B Vendors must not implement parallel agent logging mechanisms; all agent audit evidence must be retrieved exclusively via this API.

13.3 Annex C – Use-Case Definitions

13.3.1 Annex C1 – Executive list of use-cases

Code	LoB	Refer to	Name	Priority
CPCU-1	P&C Underwriting - Corporate	Candidate Vendors should also consult Attachment 1 filename "Ethniki Asfalistiki_for Automations 20251107_RFP_use_cases.pdf" for more information on nature, complexity, benefits and expected implementation timelines per each use case.	Automated Geocoding & GIS Integration	**1.2027**
CPCU-2	P&C Underwriting - Corporate		Intelligent Mailroom Automation	**1.2026**
CPCU-3	P&C Underwriting - Corporate		Information Aggregation & Pricing Automation	**1.2026**
CPCU-4	P&C Underwriting - Corporate		Public Tender Intelligence & Processing (Διαγωνισμοί)	**1.2027**
CPCU-5	P&C Underwriting - Corporate		Contract Generation & Document Automation	**1.2027**
CX-1	CX Call Centre		Call Transcription, Summarisation & CRM Population	**1.2027**
CX-2	CX Call Centre		AI Assistant to Operator - Intelligent Agent Assistant (Real-Time Support)	**1.2027**
CX-3	CX Call Centre		AI Assistant to Operator for Synallages – Transaction Automation	**1.2027**
CX-4	CX Call Centre		AI Assistant to Operator - Assessors Appointment Data Entry	**1.2027**
CX-5	CX Call Centre		Advanced Email/Form Categorization & Intelligent Routing	**1.2027**
GLH-1	Group Life & Health		Automated Offer Generation & Pricing	**1.2026**
GLH-2	Group Life & Health		Contract Generation Automation	**1.2026**
GLH-3	Group Life & Health		System Integration & Member Onboarding	**1.2026**
LHC-1	Life & Health Claims		Health Claims Email Processing & Auto-Population (Anagelies)	**1.2026**
LHC-2	Life & Health Claims		Intelligent Hospitalization Claims Routing (Nosileies)	**1.2026**
LHC-3	Life & Health Claims		Invoice Data Extraction & Validation (Timologisi)	**1.2026**
LHC-4	Life & Health Claims		AI-Powered Fraud Detection & Risk Assessment	**1.2026**
LHC-5	Life & Health Claims		Level 3 Support Query Automation	**1.2027**
LHU-1	Life & Health Underwriting		Client Entity Matching & 360° History Aggregation	**1.2026**
LHU-2	Life & Health Underwriting		Medical Documents IDP & Advisory Summaries	**1.2027**
LHU-3	Life & Health Underwriting		Additional Clauses (Prosthetes Praxeis) Mailroom Automation	**1.2027**
PCM-1	P&C Motor Claims		Automated Claim Intake (FNOL)	**1.2026**
TECH-1	Technology		IT / Service Desk Front-line and Back-office fulfilment operations	**1.2026**
TECH-2	Technology		IT / Software Development / DevOps / SecDevOps	**1.2026**

Candidate Vendors should also consult Attachment 1 filename **"Ethniki Asfalistiki_for Automations 20251107_RFP_use_cases.v.3.pdf"** for more information on nature, complexity, benefits and expected implementation timelines per each use case.

13.3.2 Annex C2 – Detailed list of use-cases

Code	LoB	Name	Description	Priority
CPCU-1	P&C Underwriting - Corporate	Automated Geocoding & GIS Integration	Batch geolocalizing service to transform property addresses from contracts (Word / Excel attachments) into XY coordinates - Automated extraction from ~3,000 existing complex multi-location contracts - Integration with GEOL GIS platform for bulk upload CSV / Excel exports	**1.2027**
CPCU-2	P&C Underwriting - Corporate	Intelligent Mailroom Automation	AI-powered classification of 5,000 daily emails across shared mailboxes. - NLP classification into Offer / Issuance / Addendum / Question - OCR in email attachment - Smart routing to responsible department / underwriter, set priority tagging, duplicate email detection - Potential enhancement to extract summaries and tasks	**1.2026**
CPCU-3	P&C Underwriting - Corporate	Information Aggregation & Pricing Automation	AI tool that ingests insurance brokers' / agents' needs assessment from clients (emails / PDFs), automatically extracts relevant data and pre-populates standardized underwriting offer templates. - Underwriters maintain full control over critical decisions (exposure assessment, re-insurance, pricing) - Integrates with multiple systems for historical data if needed	**1.2026**
CPCU-4	P&C Underwriting - Corporate	Public Tender Intelligence & Processing (Διαγωνισμοί)	Automated scanning of tender platforms for new opportunities - AI extraction of "Technical Requirements" sections from 60-70 page Greek documents (no translation needed) - Handles ~750 annual tenders - Orchestrates multi-department coordination for pricing (Property / Auto / Life / Marine) - Overall compliance versus tender requirements check	**1.2027**
CPCU-5	P&C Underwriting - Corporate	Contract Generation & Document Automation	Template-based, automated offer and contract generation - Dynamic content population from approved proposals, Greek / English language support, clause library management, compliance check, document version control - Focus on standardized packages with some customization capability	**1.2027**
CX-1	CX Call Centre	Call Transcription, Summarization & CRM Population	Speech-to-text, auto-summarization, key data extraction (customer, reason, outcome) and optional automatic Salesforce field population; includes sentiment analysis.	**1.2027**
CX-2	CX Call Centre	AI Assistant to Operator - Intelligent Agent Assistant (Real-Time Support)	Consolidated, real-time access to customer info from internal systems during calls; phases include data aggregation, proactive AI suggestions, natural-language search and predictive assistance.	**1.2027**
CX-3	CX Call Centre	AI Assistant to Operator for Synallages – Transaction Automation	Automate payment-related inquiries (status, due dates, codes, amounts) through AI-powered information retrieval and response - Handles ~150-180 payment-related daily calls -Integration with payment systems for real-time data access pending	**1.2027**
CX-4	CX Call Centre	AI Assistant to Operator - Assessors	Intelligent chatbot for automated assessor appointment scheduling with customers and repair garages - Integration with assessors' calendars for live availability and appointment assignment based on geolocation criteria	**1.2027**

		Appointment Data Entry	- Automated SMS / email confirmations, rescheduling capabilities - Garage network integration for direct booking	
CX-5	CX Call Centre	Advanced Email/Form Categorization & Intelligent Routing	Two-tier categorization system: (1) Forms: Initial categorization by customers (Phase 2, launching ~1 month), (2) AI: Secondary detailed categorization and intelligent routing beyond form capabilities. AI analyzes form submissions and remaining emails (~1,000/day) for: detailed sub-categorization, priority flagging, specialist assignment, action type identification, automated data extraction to Salesforce. Rationale: Forms provide basic categories but cannot include 100+ fields; AI completes granular routing.	**1.2027**
GLH-1	Group Life & Health	Automated Offer Generation & Pricing	Automate the complete offer workflow from agent request to delivered PDF. Platform-flexible solution. Includes: census file upload and processing, AI-powered pricing calculation using rate tables (pricing model – excel file), automated proposal document generation from templates with dynamic content population based on coverage selections, automated approval workflow, and email delivery automation.	**1.2026**
GLH-2	Group Life & Health	Contract Generation Automation	Auto-generate policy contracts based on accepted proposals. Agent uploads signed client form and UBO and other employer documentation needed → system extracts critical data (employer's data, start date, payment frequency, payment method) → generates final contract. Initial version includes AI signature detection on uploaded PDFs to validate signatures. Features: contract templates based on finalized proposals, dynamic content insertion, coverage terms auto-population, document storage and tracking. Future enhancement: full e-signature integration.	**1.2026**
GLH-3	Group Life & Health	System Integration & Member Onboarding	Automate policy data entry and insured member loading into Aviva system with payment verification gate. Phase 1: Automated load file generation from accepted proposals with manual underwriting check for advance payment confirmation. Phase 2: Direct Aviva API integration for coverage data (limits, percentages, deductibles) - critical for eRules and eClaims processing.	**1.2026**
LHC-1	Life & Health Claims	Health Claims Email Processing & Auto-Population (Anagelies)	AI-powered processing of 150 daily emails for health damage claims. Intelligent attachment extraction, multi-format OCR, automated data entry to C-MOD and Aviva systems, missing-document detection and automatic partner notification; human-in-the-loop for exceptions and threading via subject conventions.	**1.2026**
LHC-2	Life & Health Claims	Intelligent Hospitalization Claims Routing (Nosileies)	Use Accurate's (Triple) nightly feed to auto-assemble decision packs and pre-approve repetitive/standardized admissions (e.g., chemotherapy) via checklist/rules and plan terms; adjuster reviews and issues final approval; foundation for broader decision support across more case types. TIED to SUB-2	**1.2026**
LHC-3	Life & Health Claims	Invoice Data Extraction & Validation (Timologisi)	Advanced IDP/OCR for hospital invoice processing from physical and PDF documents. Multi-format invoice support, automated field mapping and extraction, business rule validation, feed to Coassess (Covariance - go-live Nov 2025), automated Aviva updates, proforma-to-invoice matching, smart approval routing, exception handling workflow.	**1.2026**
LHC-4	Life & Health Claims	AI-Powered Fraud Detection & Risk Assessment	AI-assisted claim summary and risk indication (high/med/low). extract signals from medical notes/diagnostics/receipts, reconcile with incident description and policy terms, auto-request missing info; human-in-the-loop for decisions.	**1.2026**
LHC-5	Life & Health Claims	Level 3 Support Query Automation	Digitize ~30 physical mail packets/day (~300 claims): high-throughput scanning, document classification, key-field extraction, and automated registration into the correct workflow/core (via RPA/API), with quality scoring and exception handling.	**1.2027**
LHU-1	Life & Health Underwriting	Client Entity Matching & 360° History Aggregation	Automated client recognition and dedup across departments; assemble a unified view of existing policies/claims including ALL contracts (active & canceled) within 5-year period with status change dates; scan Aviva "folder" notes for keywords (e.g., exclusions, loadings, exceptions, reinstatements) and surface a summary to the underwriter (Aviva +	**1.2026**

			Portal + Saperion).	
LHU-2	Life & Health Underwriting	Medical Documents IDP & Advisory Summaries	IDP/OCR and NLP to read medical attachments (labs, reports), normalize values, detect abnormalities, and produce a checklist/summary with flags and suggested additional exams; phased implementation starting with recommendations to build trust; underwriter remains decision-maker.	**1.2027**
LHU-3	Life & Health Underwriting	Additional Clauses (Prosthetes Praxeis) Mailroom Automation	Automate intake of ~120k addendum requests/year (email/physical): classify request type (key types: beneficiary changes ~800, payment method ~7,000, address ~6,500), extract key fields from forms, validate vs policy data, and update Aviva via API or RPA; attach documents to Saperion; handle exceptions; support batch processing.	**1.2027**
PCM-1	P&C Motor Claims	Automated Claim Intake (FNOL)	AI workflow to capture First Notice of Loss via portal/web/phone/email/messaging: voice AI agent, photo/video upload, document OCR, geolocation, completeness checks, basic fraud/duplicate screening, auto-creation of claim file.	**1.2026**
TECH-1	Technology	IT / Service Desk Front-line and Back-office fulfilment operations	Improve Service Desk efficiency and user experience by introducing AI-assisted capabilities across front-line support (Self-Service Portal assistant) and back-end technical teams (Contributor assistant & Autonomous agents). The solution must ideally leverage the native capabilities of the enterprise ITSM platform (ServiceNow), or seamlessly integrate to it, to automate triage, accelerate resolution, and enhance knowledge usage while ensuring full GDPR and EU AI Act compliance.	**1.2026**
TECH-2	Technology	IT / Software Development / DevOps / SecDevOps	Introduce AI-assisted capabilities across the software delivery lifecycle to improve developer productivity, accelerate release cycles, enhance code quality and security, and support large-scale refactoring and modernization of legacy applications.	**1.2026**

Candidate Vendors should also consult Attachment 1 filename **"Ethniki Asfalistiki_for Automations 20251107_RFP_use_cases.v.3.pdf"** for more information on nature, complexity, benefits and expected implementation timelines per each use case.

13.4 Annex E – Glossary

Term	Definition
AI System	A machine-based system operating with varying autonomy and adaptiveness.
High-Risk AI System	AI falling under EU AI Act high-risk classification.
Retrieval-Augmented Generation (RAG)	Architecture combining retrieval with generative AI.
Vector Store	Storage layer for embeddings enabling semantic retrieval.
Embedding	Vector representation of text used for similarity search.
LLM Gateway	Controlled routing layer for model access.
Explainability	Ability to interpret AI reasoning.
Model Drift	Performance deviation of models over time.
Human-in-the-Loop (HITL)	Mandatory human oversight.
Operational Resilience	DORA-aligned continuity requirements.
ICT Incident	DORA-defined ICT service disruption.
Personal Data	GDPR-defined identifiable information.
Data Controller / Processor	GDPR-defined processing roles.
DPIA	Impact assessment for high-risk processing.
TOMs	GDPR Article 32 security measures.
Sub-processor	Third party processing on Vendor's behalf.
Shadow IT	Unapproved infrastructure or models.
Onboarding	Integration into LOT A Control Plane.
SLA	Contractual service performance metrics.
Commercially Available AI Platform (CAAP)	A commercially available, purpose-built AI platform proposed by a LOT B Vendor as the primary delivery solution for a use-case work package, in place of a custom-built integration with the LOT A Control Plane. A CAAP must obtain Certified External Platform (CEP) status before promotion to UAT.
Certified External Platform (CEP)	A designation granted by the LOT A Vendor (on behalf of Ethniki) to a Commercially Available AI Platform that has satisfied all eligibility criteria and mandatory integration obligations defined in Section 7.9. CEP status permits a LOT B Vendor to deploy the platform's embedded AI components under a structured governance bridge, without categorical exclusion from the procurement.
CEP Register	A maintained record, managed by the LOT A Vendor on behalf of Ethniki, listing all platforms that have achieved, failed, or had Certified External Platform (CEP) status revoked. Subject to quarterly compliance review.
CEP Governance Bridge	The set of mandatory technical integrations that a Certified External Platform (CEP) must implement to bring its AI operations within the audit, observability, and compliance authority of the LOT A Control Plane. Defined in Section 7.9.2. Includes audit log export, LLM Gateway compliance or bridge, data residency attestation, DORA dependency mapping, EU AI

	Act risk classification, Dynatrace observability integration, and incident/change notification.
Agentic AI	An AI system design pattern in which one or more AI models autonomously execute multi-step workflows by selecting and invoking external tools, APIs, or sub-agents in sequence, making decisions based on intermediate results, and operating with a perceive–reason–act loop until a goal condition is reached or a HITL checkpoint is triggered. Distinct from standard LLM completion or RAG patterns, which are stateless and single-turn.
Agent Orchestration Layer	The platform component responsible for managing the lifecycle, execution, state, identity, and governance of AI agents operating within the enterprise AI environment. Defined as the optional Component H in Section 6.3 of this RFP. Where Component H is deployed, it serves as the sole permitted runtime for agentic workflows within the LOT A Control Plane boundary.
Tool Registry	A governed catalogue of tools, APIs, and external integrations that AI agents are permitted to invoke within the LOT A Control Plane. Each registered tool carries mandatory metadata including owning system, data classification, permitted callers, and approval status. Agents may only invoke tools explicitly listed in the registry with scoped caller permission; unapproved invocations are blocked and logged. Part of Component H (Section 6.3) where proposed.
HITL (as Orchestration Primitive)	In the context of agentic workflows (Component H), Human-in-the-Loop (HITL) is elevated beyond exception handling to a first-class orchestration primitive: a structured interrupt-and-resume mechanism by which an agent pauses execution at a defined checkpoint, routes a decision request to a designated human reviewer with full context, and resumes processing with the reviewer's input incorporated. Distinct from the standard HITL exception handling defined elsewhere in this RFP, which applies to non-agentic AI-assisted decisions.

13.5 Annex F – Compliance Framework (GDPR, EU AI Act, DORA)

13.5.1 Annex F1 – GDPR Compliance Pack

This Annex defines the minimum set of GDPR artefacts and documentation that each Vendor must provide and keep up to date for the duration of the Contract. At a minimum, the GDPR Compliance Pack shall include:

- Draft Data Processing Agreement (DPA) aligned with Article 28 GDPR and Ethniki's standard clauses, including detailed allocation of responsibilities and incident notification timelines.
- A register of sub-processors used or intended to be used for the provision of the Services, including their location, role and the technical and organisational measures they implement.
- A Technical and Organizational Measures (TOMs) Catalogue describing, in sufficient detail, the controls in place for access management, encryption, network and infrastructure security, physical security, secure development, logging and monitoring, backup and recovery, and secure disposal of data and assets.
- Support materials for Data Protection Impact Assessments (DPIAs) for the relevant processing activities, including description of purposes, categories of data subjects and data, processing operations, risk analysis and mitigation measures.
- Data flow diagrams or equivalent documentation showing the flow of personal data between components, systems and locations, including any transfers outside the EEA, and the associated retention periods.

13.5.2 Annex F2 – EU AI Act Compliance Pack

This Annex specifies the documentation expected from Vendors to support EU AI Act-aligned governance of AI components used in the Services. The EU AI Act Compliance Pack shall include, as applicable:

- The Vendor shall include an AI Risk Management Plan covering risk identification, mitigation measures, monitoring controls, and residual risk analysis across the AI lifecycle.
- A Data Governance Specification describing the datasets used for training, validation and testing, including their sources, representativeness, known limitations, pre-processing steps, labelling practices and bias-mitigation activities.
- For all datasets used in training, validation, or fine-tuning, the Vendor shall provide Dataset Data Sheets documenting provenance, representativeness, data quality indicators, governance controls, and known limitations.

- Technical Documentation of the AI system(s), including high-level architecture diagrams, model types and parameters, training procedures, evaluation methodology and key performance metrics (accuracy, robustness, coverage, etc.).
- Logging and Traceability Specifications outline which events are logged, how long logs are retained, how they are protected, and how they can be queried for audit and investigation purposes.
- A Human Oversight Plan describing the role of human operators in supervising AI outputs, the escalation paths, and the circumstances in which human review or override is mandatory.
- Where applicable, evidence packages or declarations relating to conformity assessment activities performed under the EU AI Act or applicable sectoral guidance.

13.5.3 Annex F3 – DORA ICT Risk & Operational Resilience Pack

This Annex sets out the minimum documentation required from Vendors to support Ethniki's obligations under DORA in relation to ICT risk management and operational resilience. The DORA ICT Risk & Operational Resilience Pack shall include:

- An ICT Risk Register specific to the Services delivered, outlining key risks, their likelihood and impact, and the controls implemented to mitigate them.
- Business Continuity and Disaster Recovery documentation for the critical components supporting the Services, including recovery objectives, dependency mappings and fallback procedures.
- An ICT Incident Management procedure describing detection, classification, escalation, containment, eradication and recovery activities, as well as communication with Ethniki and relevant third parties.
- A description of the Vendor's third-party risk-management practices, including how critical suppliers are identified, assessed and monitored, and how concentration risk is managed.
- Exit, data portability and transition plans that would apply in the event of contract termination, insourcing or transition to another provider, including timelines, responsibilities and data-handover mechanisms.
- Evidence or templates relating to resilience testing (e.g. DR tests, fail-over tests, cyber-range exercises) relevant to the Services, including, at a minimum, test objectives, scenarios, frequency and reporting format.

13.6 Annex H – Architecture & Integration Requirements

13.6.1 High-Level Enterprise AI/RAG Control Plane Architecture

- The Control Plane shall provide a unified enterprise foundation comprising:
- Ingestion & Pre-Processing Layer: OCR/ICR, parsing, metadata extraction, semantic chunking, image-to-text, PDF decompositions, email/thread flattening.
- Embedding & Vectorization Layer: Embedding models, re-embedding workflows, quality scoring, semantic consistency checks.
- Vector Store: Scalable dense/sparse/hybrid indexes, metadata filtering, RBAC propagation, TTL-retention rules, lineage tracking.
- Retrieval Orchestration: Hybrid search, reranking, contextual filtering, safety-aware retrieval, query rewriting, multi-hop retrieval.
- LLM Gateway: Policy-driven routing, temperature control, cost accounting, sensitive-data redaction, prompt-shielding, offline/batch inference support.
- Prompt Management: Versioning, template inheritance, guardrails, parameterized RAG prompts, A/B testing framework.
- Governance Layer: Logging, auditability, explainability, drift detection, human-in-the-loop escalation, KRI dashboards.
- Observability & AI-Ops: Token usage analytics, retrieval quality scores, anomaly detection, environment health, SLO/SLA tracking.

13.6.2 Integration Principles for LOT Bx Vendors

- API-First: All interactions with ingestion, retrieval, generation, and metadata services occur via approved Control Plane APIs.
- No Local AI Components: Vendors may not deploy private vector DBs, model runtimes, prompt engines, or retrieval pipelines, except where those components form part of a Certified External Platform (CEP) with active CEP status under Section 7.9. CEP-designated platforms must satisfy all seven integration obligations in Section 7.9.2 and remain subject to LOT A audit and certification authority.
- Data Flow Constraints: All data entering or exiting LOT Bx solutions must pass through the Control Plane for governance and audit.
- Security Enforcement: IAM/SSO, RBAC, logging, encryption, and SIEM hooks must be inherited from LOT A.
- Model Interoperability: Vendors must use models exposed by the gateway; external models must not be called directly.
- Deployment Standards: CI/CD pipelines, artifact scanning, environment segregation, and release approvals follow enterprise templates.

13.6.3 Environment Model & Operational Architecture

- Ethniki requires clear and isolated environments:
- DEV: Developer sandboxes, mock data, experiment tracking, automated linting and security scanning.
- TEST: Automated functional, integration, and performance tests; DPIA/AI-risk pre-validation.
- PRE-PROD: Production-like configuration, performance baselines, resilience testing, rollback trials.
- PROD: Fully audited runtime with strict change control, monitoring, human-supervision triggers, and incident response processes.

All environments must include:

- Immutable infrastructure or containerization standards.
- Automated promotion workflows with approvals.
- Environment-specific encryption keys and IAM/SSO policies.

13.6.4 Security Architecture Requirements (Extended)

- Zero-Trust Enforcement: Network micro-segmentation, continuous verification, adaptive policies.
- IAM Integration: SSO, SCIM provisioning, per-use-case RBAC scopes, just-in-time access.
- Encryption: support for cryptographic modules, envelope encryption for sensitive retrieval metadata.
- Auditability: Tamper-evident logs, forensic export capability, regulator-ready evidence packages.
- SIEM Integration: Real-time forwarding of auth, retrieval, model invocation, exception and anomaly events.
- DORA Alignment: RTO/RPO targets, failover zones, cross-region replication, periodic resilience drills.
- The Vendor must perform robustness evaluation against prompt injection, adversarial manipulation, unsafe completions, and model override attempts. Evidence of safety guardrail validation shall be included in acceptance testing.

13.6.5 Data Flow, Lifecycle & Governance Framework

- Data Acquisition: All ingestion must apply GDPR-aligned minimization and classification rules.
- RAG Lifecycle: Version control for embeddings, indexes, prompts, and retrieval logic; automated lineage mapping.
- HITL Branching: Human validation steps for high-risk processes, configurable by LoB.
- Model Governance: Approval workflows, dataset documentation, validation metrics, fairness/bias checks. The Vendor shall conduct fairness and bias assessment for all AI

components impacting underwriting, claims assessment, or customer interaction flows. A bias testing report must be submitted as part of acceptance criteria.

- Drift Monitoring: Statistical drift detection on embeddings, model weights, and retrieval patterns.
- Change Control: Mandatory impact assessment, rollback planning, and validation before promoting updates.

13.6.6 Vendor Onboarding & Compliance Sequence

- Architectural alignment workshop.
- Security & compliance review (GDPR/DORA/AI Act).
- API & data-flow design validation with LOT A team.
- Ingestion pipeline configuration and sample dataset tests.
- Retrieval & RAG prompt calibration.
- Explainability & HITL workflow confirmation.
- CI/CD and deployment readiness assessment.
- Go-Live approval with audit artefacts.
- Postproduction monitoring setup and reporting routines.

13.6.7 Performance, Scalability & Capacity Requirements

- Horizontal scaling of ingestion, vector store, and retrieval nodes.
- Latency SLOs: <1.0 sec retrieval, <2.5 sec RAG response under peak load.
- Throughput: Support for concurrent workloads across all LOT B use cases.
- Capacity Planning: Automated autoscaling, usage forecasting dashboards.

13.6.8 Logging, Traceability & Explainability Framework

- Mandatory Logging: Inputs, retrieved chunks, model metadata, outputs, errors, overrides.
- Traceability: Ability to reconstruct the entire chain from request to generated answer.
- Explainability: Token attribution, retrieval contribution ranking, explainers.
- Evidence Packs: Vendor must deliver regulator-ready explainability artefacts.

13.6.9 Prohibited Architectural Patterns

- Direct calls to external LLM APIs, except where the calling system holds active Certified External Platform (CEP) status under Section 7.9 and routes all inference events to the LOT A audit log via the CEP governance bridge.
- Vendor-side embedding stores or retrieval layers, except within a CEP-designated platform operating under the integration obligations of Section 7.9.
- Local model hosting without approval.
- Bypassing Control Plane logs, SIEM, or prompt governance.
- Shadow databases, hidden caches, or undocumented enrichment layers.

13.7 Annex I – Vendor Compliance Checklists (LOT A & LOT B)

The purpose of this Annex is to ensure that all Bidders provide complete, consistent, and comparable information. All responses must be submitted using the templates provided in this Annex. Responses that do not follow the required structure may be deemed non-compliant.

13.7.1 Annex I.1 - General instructions applicable to all LOTs

- Complete All Applicable Sections:
    - Bidders must complete all lines and all fields in the response template applicable to their LOT.
    - LOT A Bidders complete 13.8.4
    - LOT B Bidders complete 13.8.5
- Do Not Modify Structure:
    - The format, table structure, numbering, and headers must not be altered.
    - Any modifications may lead to disqualification.
- Evidence Must Be Verifiable:
    - For each requirement, Bidders must clearly reference the supporting document, slide deck, appendix, or URL submitted in their proposal (e.g., “Appendix 4, p.12”).
    - Vague references (“explained in technical section”) are insufficient.
- Clarity and Specificity Required:
    - Responses must be specific, measurable, and unambiguous. Marketing language or generic statements are not accepted without evidence.
- All Mandatory Requirements Are Binding:
    - Any mandatory requirement marked as Non-Compliant or Partially Compliant may disqualify the Bid.
    - Where “Partial” is selected, a remediation plan and timeline must be provided, or otherwise explanation as to why does not weaken the proposal against the baselines set by the RFP.
- Terminology Alignment:
    - All terminology must align with the definitions and terminology used throughout the RFP (Sections 6, 7, 8–10, Annex A, Annex H, etc.).

- Non-Compliance Protocol - If the Bidder cannot comply fully with a requirement:
  - They must select Partial or No.
  - Provide justification.
  - Provide remediation, alternatives, or workaround where applicable.
  - Failure to provide justification may result in rejection.
- Confidentiality and IP
  - All submitted material must clearly indicate IP ownership.
  - Material not identified as proprietary may be considered open for evaluation and storage.

13.7.2 Annex I.2 - Instructions for LOT A Bidders (Control Plane Vendor)

- LOT A Bidders are responsible for delivering the enterprise AI/RAG Control Plane. Responses must:
  - Address every requirement listed under the LOT A template (A1.1–A7.2).
  - Demonstrate a complete and integrated architecture covering:
  - Retrieval and vector store
  - Embeddings pipeline
  - LLM gateway and routing
  - Prompt governance
  - Observability and AI-Ops
  - Security, GDPR, DORA, EU AI Act compliance
- Provide evidence for:
  - Architecture diagrams
  - Monitoring dashboards
  - Operational runbooks
  - Security documentation
  - Vendor onboarding kits
  - Be consistent with the delivery expectations in Section 13 and Annex K.

Important:

- Statements such as “Our standard platform provides this capability” are not sufficient.
- Each requirement must be explicitly addressed in the template.

13.7.3 Annex I.3 - Instructions for LOT B Bidders (Use-Case Vendors)

- LOT B Bidders are responsible for delivering functional use-cases and must integrate with the LOT A platform. Responses must:
    - Describe the end-to-end functional workflow including HITL, exceptions, routing, business rules.
    - Demonstrate correct integration with:
    - LOT A’s retrieval APIs
    - Prompt governance
    - Logging and traceability
    - LLM gateway controls
    - Provide specific prompts, retrieval strategies, and grounding methods used in the use-case.
- Submit compliance mappings for:
    - GDPR data processing
    - EU AI Act classification (risk category)
    - DORA ICT dependencies
    - Provide clear UAT methodology, acceptance criteria, and performance KPIs for the use-case.
    - Include user journeys, screens, or interaction flows where relevant.

Important:

- LOT B Bidders must not describe or propose platform components already in LOT A’s scope (e.g., vector database, LLM gateway, ingestion pipelines) as part of their own solution architecture. Any such proposals will be disregarded for evaluation purposes. Exception: where a LOT B Bidder proposes a Commercially Available AI Platform (CAAP) under the CEP pathway (Section 7.9), the Bidder must fully describe the CAAP’s embedded AI components as part of the mandatory CEP disclosure in checklist items B7.1 and B7.2. Such descriptions are required for CEP certification and will not be treated as a prohibited proposal of LOT A components.

13.7.4 Annex I.4 - LOT A Vendor Compliance Checklist

It is mandatory for Vendors' Technical offers to include an Excel version of the following compliance checklist.

No.	Requirement Category	Detailed Requirement	Vendor Response (yes/partial/no)	Evidence Reference	Section Reference
A1.1	Control Plane Architecture	Provide complete high-level architecture diagram, including retrieval, embedding, gateway, observability, governance layers.			Sec 6, Annex A
A1.2	Control Plane Architecture	Provide low-level diagrams for ingestion pipelines, embedding generation, chunking logic, index management, and re-embedding workflows.			Sec 6.1(A–B)
A1.3	Control Plane Architecture	Describe storage, compute, scaling and HA/DR architecture for vector DB, gateway, and AI services.			Annex H
A2.1	Retrieval & Vector Store	Describe retrieval architecture including hybrid (dense/sparse) retrieval, RBAC filtering, metadata filtering, and index lifecycle management.			Sec 6.1(A)
A2.2	Retrieval & Vector Store	Provide retrieval quality evaluation metrics, tools, and methodologies.			Annex H
A3.1	LLM Gateway	Describe routing strategies (cost, latency, sensitivity) and supported LLM families/models.			Sec 6.1(C)
A3.2	LLM Gateway	Describe grounding enforcement, safety layers, and hallucination mitigation methods.			Annex A
A4.1	Prompt Governance	Provide prompt library structure, version control workflows, approval flows, and guardrails.			Sec 6.1(D)
A4.2	Prompt Governance	Describe parameterisation and inheritance models for prompts.			Annex A
A5.1	Monitoring & AI-Ops	Provide monitoring dashboards, drift detection methods, latency/error metrics, anomaly detection.			Sec 6.1(F)
A5.2	Monitoring & AI-Ops	Describe integration with enterprise SIEM and alerting flows.			Annex H

A6.1	Security & Compliance	Provide GDPR controls, retention, audit logging, deletion processes.			Sec 9, Annex F
A6.2	Security & Compliance	Provide EU AI Act documentation including risk classification, transparency, and logging obligations.			Sec 10, Annex F2
A6.3	Security & Compliance	Provide DORA ICT controls including DR, continuity, incident response.			Sec 8, Annex F3
A7.1	Vendor Onboarding Kits	Provide APIs, SDKs, schema definitions, sandbox environments.			Annex B, Annex A(A.3)
A7.2	Vendor Onboarding Kits	Provide integration guide and operational contract with LOT B vendors.			Annex L

13.7.5 Annex I.5 - LOT B Vendor Compliance Checklist

No.	Requirement Category	Detailed Requirement	Vendor Response (yes/partial/no)	Evidence Reference	Section Reference
B1.1	Functional Use Case Delivery	Describe full workflow, business rules, exceptions, HITL (Human-in-the-Loop) logic.			Section 7
B1.2	Functional Use Case Delivery	Provide UI/UX mockups or flow diagrams for the automation.			Annex L
B2.1	Integration with LOT A	Describe how your solution calls the LOT A retrieval API and how grounding is used.			Sec 6, Annex B
B2.2	Integration with LOT A	Describe how you consume prompts and follow governance constraints from LOT A.			Sec 6.1(D), Annex A
B2.3	Integration with LOT A	Explain how logs, traceability and auditability align with Control Plane.			Annex F
B3.1	Prompt & Retrieval Configuration	Provide example prompts and grounding logic for your use case.			Sec 6.1(D), Annex A
B3.2	Prompt & Retrieval	Provide retrieval performance expectations and fallback logic.			Annex H

	Configuration				
B4.1	Compliance	Provide GDPR mapping for personal data processed within the use-case.			Sec 9, Annex F1
B4.2	Compliance	Provide EU AI Act classification for this use-case and documentation obligations.			Sec 10, Annex F2
B4.3	Compliance	Provide DORA dependency mapping on LOT A platform and resilience expectations.			Sec 8, Annex F3
B5.1	Operational Readiness	Provide user guides, UAT scripts, acceptance criteria, support model.			Annex L
B5.2	Operational Readiness	Describe incident flows and vendor support model for your specific use-case.			Section 13
B6.1	Quality & Performance	Provide KPIs, SLAs, accuracy measures, and reliability metrics for the automations.			Section 11.B
B6.2	Quality & Performance	Describe how drift, errors, hallucinations will be handled at the use-case layer.			Annex H
B7.1	CEP Disclosure (if applicable)	If proposing a Commercially Available AI Platform (CAAP) as the primary solution for this LOT, disclose: (a) intent to seek CEP status; (b) platform name and vendor; (c) preliminary CEP compliance gap assessment against the obligations in Section 7.9.2. LOT B vendors retain full contractual accountability for CEP platform performance.			Sec 7.9
B7.2	CEP Integration Evidence (if applicable)	For CEP-designated platforms: confirm acceptance of all seven mandatory integration obligations (Section 7.9.2): audit log export, LLM Gateway compliance or bridge, data residency attestation, DORA dependency mapping, EU AI Act risk classification, Dynatrace observability integration, and incident/change notification. Provide evidence of production deployment at minimum two regulated European insurance organisations.			Sec 7.9.1, 7.9.2

Note, with respect to any references to Section 6, the requirement is “must consume, not must deliver.”. LOT B vendors must consume these platform capabilities delivered by LOT A.

13.8 Annex J: Ethniki Insurance Technology Landscape

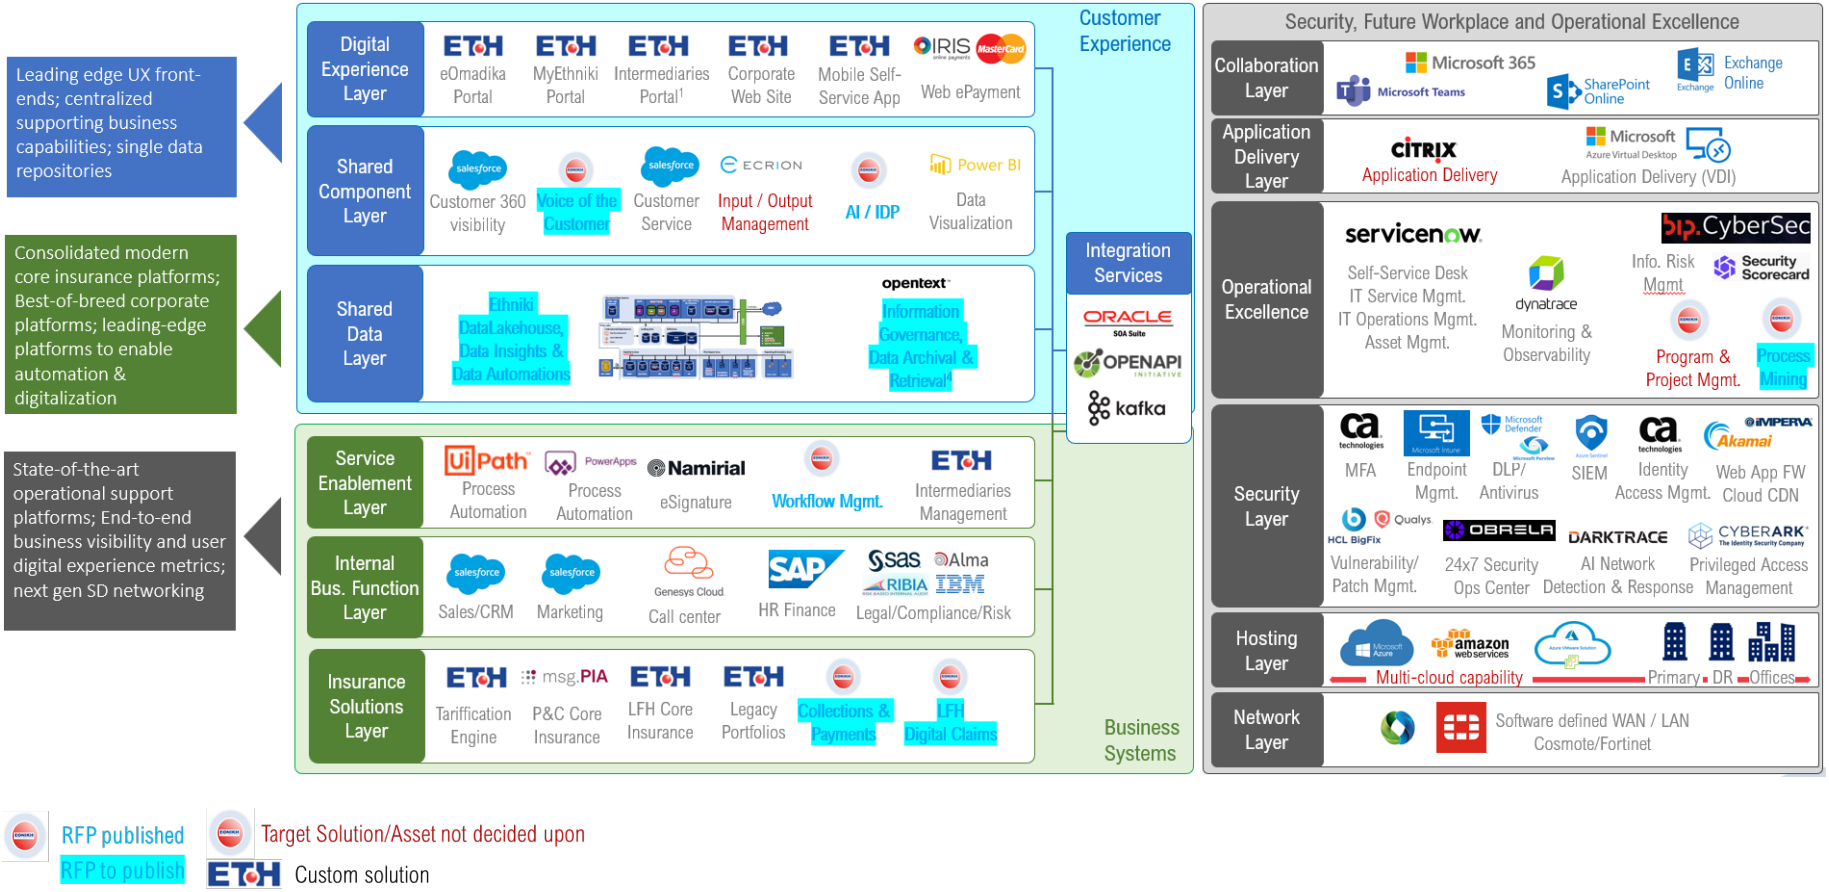

13.9 Annex K – Use-Case Delivery Playbook (Tentative Framework for Bidder Consideration)

This Playbook outlines a tentative delivery framework intended to guide LOT B bidders in preparing their proposals for the design and implementation of AI/IDP use cases onboarding to the Enterprise AI/RAG Control Plane.

It provides a baseline delivery approach reflecting Ethniki's expectations for structure, quality, compliance, and integration discipline, while allowing bidders to propose their own methodology provided all mandatory requirements are fully satisfied.

Mandatory vs. Advisory Elements:

- MANDATORY: Requirements that all LOT B Vendors must follow.
- ADVISORY: Recommended steps that Vendors may adapt as long as mandatory obligations remain satisfied.

13.9.1 Delivery Phases

- Phase 0 – Mobilization (MANDATORY) - Architecture alignment with LOT A, onboarding checklist, API conformity, security/GDPR/AI Act/DORA briefing.
- Phase 1 – Discovery (MANDATORY) - Process mapping, data minimization, exception/HITL logic, initial compliance scoping. (ADVISORY: Additional business analysis.)
- Phase 2 – Design (MANDATORY) - Functional Specification (FS), Technical Specification (TS), retrieval/prompt configuration (if RAG), GDPR mapping, AI Act classification, DORA dependency mapping. (ADVISORY: UX mockups.)
- Phase 3 – Build & Integration (MANDATORY) - Use of LOT A APIs (retrieval, embeddings, gateway, prompts, audit), logging/traceability, observability hooks, CI/CD deployment readiness. (ADVISORY: Vendor DevOps accelerators.)
- Phase 4 – Testing (MANDATORY) - Functional, integration, performance, security and compliance testing. Retrieval accuracy & grounding tests (if RAG). Drift / hallucination / bias checks where applicable. (ADVISORY: Additional methodologies.)
- Phase 5 – Deployment (MANDATORY) - Runbooks, UAT sign-off, updated DPIA/FRIA, go-live approval.
- Phase 6 – Hypercare (MANDATORY) - Monitoring, drift checks, issue resolution, acceptance package.

13.10 Annex L – LOT A Control Plane Delivery Playbook (Mandatory)

This Playbook defines the **mandatory delivery approach** for the LOT A Vendor responsible for the Enterprise AI/RAG Control Plane. Unlike LOT B Vendors, who deliver individual use cases, the LOT A Vendor delivers the **shared platform** that underpins, governs, and enables all AI capabilities across Ethniki.

LOT A has two equally critical responsibilities:

- **Design, build and operate** the AI/RAG Control Plane; and
- **Enable, support and govern all LOT B Vendors**, ensuring they can efficiently design, build, test, deploy and monitor their solutions in full alignment with platform architecture, compliance obligations, and quality standards.

The success of every LOT B use case depends on the capability, stability, documentation, and support provided by LOT A. Accordingly, all steps and obligations outlined below are **MANDATORY**.

13.10.1 Delivery Phases

- Phase 0 – Mobilization & Onboarding (MANDATORY)
  - Establish joint governance with Ethniki and initialize the Control Plane roadmap.
  - Provision environments, credentials, documentation and API references for LOT B Vendors.
  - Conduct onboarding sessions on architecture, API usage, compliance (GDPR, EU AI Act, DORA), and development standards.
- Phase 1 – Platform Architecture & Design (MANDATORY)
  - Design ingestion pipelines, embedding workflows, retrieval layer, LLM gateway, prompt governance, auditability and observability components.
  - Define multi-vendor onboarding standards, reference patterns, sample prompts, SDKs, test harnesses and certification criteria.
  - Produce architecture, security and compliance design artefacts.
- Phase 2 – Build & Enablement (MANDATORY)
  - Implement all Control Plane components.
  - Deliver sandbox environments, SDKs, CLI tools, templates and example flows to accelerate LOT B development.
  - Provide LOT B Vendors with standardized scaffolding for retrieval evaluation, prompt versioning, logging and monitoring.
- Phase 3 – Integration Support & Multi-Vendor Interaction (MANDATORY)
  - Support each LOT B Vendor through workshops, reviews and technical guidance.
  - Validate each vendor's use of retrieval, prompts, embeddings, gateway routing, traceability and observability.
  - Assist with compliance activities (DPIA inputs, AI Act classification, DORA dependencies).
- Phase 4 – Testing & Certification (MANDATORY)

    - Provide LOT B Vendors with certification tests for retrieval accuracy, latency, grounding quality, logging completeness and explainability.
    - Assist in troubleshooting, tuning and resolving integration defects.
- Phase 5 – Go-Live Coordination & Operations (MANDATORY)
    - Coordinate releases across multiple LOT B Vendors.
    - Provide dashboards, SRE processes, incident handling and change-control mechanisms.
    - Ensure operational stability, continuity and regulatory evidence retention.
- Phase 6 – Continuous Enablement & Evolution (MANDATORY)
    - Provide ongoing guidance, retraining workflows, drift monitoring and optimization support.
    - Lead platform scalability, risk management and continuous improvement.

13.11 Annex M – Group AI Infrastructure: Strategic Context and Evaluative Preference

This Annex provides contextual information on Ethniki Asfalistiki's strategic interest in Group AI infrastructure alignment, as referenced in Section 6.6. It does not form part of the mandatory technical specification and creates no obligation on vendors. It is provided to help LOT A bidders understand Ethniki's strategic direction of travel and to encourage well-informed proposals from vendors with relevant experience.

13.11.1 Background

Ethniki Asfalistiki is a member of the Piraeus Bank Group. Piraeus Bank has been the first Greek financial institution to publicly communicate an enterprise AI strategy and has invested significantly in building a production-grade AI infrastructure programme. The Bank's Data & AI Hub, operating under a Hub and Spoke governance model, has identified and is delivering over 70 AI use cases across its business functions over a four-year programme horizon (2025–2028).

Central to this programme is an enterprise AI Control Plane built on Microsoft Azure AI Foundry, which Piraeus has established as the shared technical foundation for all AI use cases across the organisation. This infrastructure is operational and has been designed and validated against the same regulatory environment – GDPR, EU AI Act, and DORA – that applies to Ethniki.

13.11.2 The Piraeus Group AI Foundry Architecture

The Piraeus Group AI Foundry architecture addresses all mandatory LOT A platform components specified in Section 6.2 of this RFP. The table below summarises the current implementation status as understood by Ethniki at the time of this RFP's publication. Vendors are advised that this information reflects Ethniki's current understanding and that formal due diligence on the Piraeus infrastructure is ongoing. Vendors should not treat this table as a definitive technical specification.

LLM Gateway & Model Routing: Azure AI Foundry model catalogue with routing across multiple frontier models including Anthropic Claude, OpenAI GPT-4o, and Microsoft Phi-3; token-based cost allocation and model-level observability. Status: Production.

Vector Store & Retrieval (RAG): Azure AI Search with hybrid semantic and keyword retrieval, integrated with Azure Blob and SharePoint content sources. Status: Production.

Document Ingestion & IDP Pipeline: Azure Document Intelligence for structured extraction; multi-format ingestion pipeline covering PDF, DOCX, email, and image formats with classification routing. Status: Production.

Prompt Management & Versioning: Azure AI Foundry Prompt Flow with version-controlled prompt templates, A/B testing, and evaluation harness. Status: Production.

Observability & AI-Ops: Dynatrace integration for infrastructure-level monitoring (consistent with Section 6.2(F) of this RFP); Azure Monitor for AI-specific telemetry and unified dashboarding. Status: Production.

Security & Data Residency: All compute within Azure West Europe region; Azure Private Link for network isolation; Azure Key Vault for secrets management; EU data residency attestation in place. Status: Production.

GDPR & Audit Logging: Centralised audit log in Azure Log Analytics Workspace; data classification tagging; PII redaction pipeline; GDPR-compliant data processor agreements with all sub-processors. Status: Production.

CI/CD & Environment Management: Azure DevOps pipelines; DEV/UAT/PROD environment promotion gates with automated compliance checks and human approval gates. Status: Production.

Agent Orchestration & Execution Layer (equivalent to Section 6.3 optional component): LangGraph-based agent runtime with MCP (Model Context Protocol) for tool integration and agent identity management via Entra ID Managed Identities. Status: Roadmap (H2 2026).

13.11.3 Why Group Alignment Is Strategically Attractive to Ethniki

Regulatory and compliance alignment. Piraeus Bank and Ethniki operate within overlapping regulatory environments: both are subject to GDPR, the EU AI Act, and DORA. The Piraeus AI Foundry architecture has been designed and validated within these constraints. Ethniki could leverage Piraeus's existing compliance artefacts – DPIAs, EU AI Act risk assessments, DORA ICT third-party risk registers, and Dynatrace observability runbooks – as starting templates for equivalent Ethniki use cases, reducing the compliance build effort materially.

Production-validated blueprint. The Piraeus AI Control Plane is not a reference architecture on paper; it is operational infrastructure supporting live use cases including customer communication AI, credit risk document processing, and regulatory compliance automation. An LOT A vendor building on this blueprint would be extending a proven platform rather than constructing one from a greenfield specification, which Ethniki considers a meaningful risk-reduction factor.

Group-wide economies of scale. Converging on a shared Group infrastructure standard creates the potential for: consolidated enterprise vendor relationships covering both Piraeus and Ethniki; shared

security operations and SOC/SIEM frameworks that can be extended to the Ethniki AI environment; and talent and knowledge transfer from the Piraeus Digital, Data & AI division to the Ethniki LOT A vendor team.

13.11.4 Status of Due Diligence and Implications for This Procurement

Ethniki's due diligence on the Piraeus Group AI infrastructure is ongoing.

This procurement therefore proceeds on a technology-agnostic basis. The information in this Annex is provided as strategic context only. Ethniki reserves the right to issue a further update to this RFP, or to reflect Group infrastructure alignment requirements in a subsequent procurement, once due diligence is complete. LOT A vendors are not required to propose Azure AI Foundry, and proposals based on other platforms will be evaluated on equal terms against the mandatory criteria in Sections 6.2, 6.3, and 11.1.

13.11.5 Guidance for LOT A Vendors

LOT A vendors wishing to articulate Group infrastructure alignment in their proposal are encouraged to address the following questions within the Vendor Innovation & Delivery Model section of their technical submission:

- Describe your organisation's experience implementing and operating Microsoft Azure AI Foundry at enterprise scale, with specific reference to the mandatory components in Section 6.2 of this RFP.
- Describe how your proposed LOT A solution could achieve technical interoperability or infrastructure alignment with an existing Azure AI Foundry-based Group AI platform, and what the practical implications would be for governance, data flows, security, and compliance.
- Describe any existing relationships with Piraeus Bank Group or Microsoft as relevant to this engagement, including any Group enterprise agreements that could be extended to cover Ethniki's requirements.

Responses to these questions are optional and will not disadvantage vendors who do not provide them. They will be taken into account as a positive differentiator where relevant and credible.

13.12 Annex N – Q&A Submission Template

Vendors must use this template for all clarification questions submitted during the Q&A period. Each question must occupy a separate row. Submit the completed template to the designated Procurement contact address provided in the covering communication. Incomplete submissions — missing mandatory fields or bundling multiple questions in a single row — will be returned without a response. Column widths are indicative; vendors may adjust formatting for readability.

#	LOT Reference	Section Reference	Use Case Ref (if applicable)	Question Type	Requirement Text Referenced (verbatim)	Question	Proposed Interpretation (optional)
1	e.g. LOT B4	e.g. Section 7.9.2	e.g. LHC-3	e.g. Technical / Architectural	"The platform vendor accepts the technical integration requirements defined in Section 7.9.2…"	Please confirm whether the LLM Gateway compliance requirement in Section 7.9.2 permits a bridging integration as an alternative to direct gateway routing in all cases, or only where a technical exemption has been granted.	We interpret "or demonstrate equivalent governance via a LOT A-approved bridging integration" as permitting a bridge in all cases where direct gateway routing is architecturally impractical, subject to LOT A Vendor approval.
2							
3							

Question Type options: Technical / Architectural | Commercial / Pricing | Legal / Contractual | Compliance / Regulatory | Process / Procedural

LOT Reference options: LOT A | LOT B1 | LOT B2 | LOT B3 | LOT B4 | LOT B5 | LOT B6 | LOT B7 | General / All LOTs